

# CPSC 4040/6040

# Computer Graphics

# Images

Joshua Levine  
[levinej@clemson.edu](mailto:levinej@clemson.edu)

# Lecture 12

# Region Processing

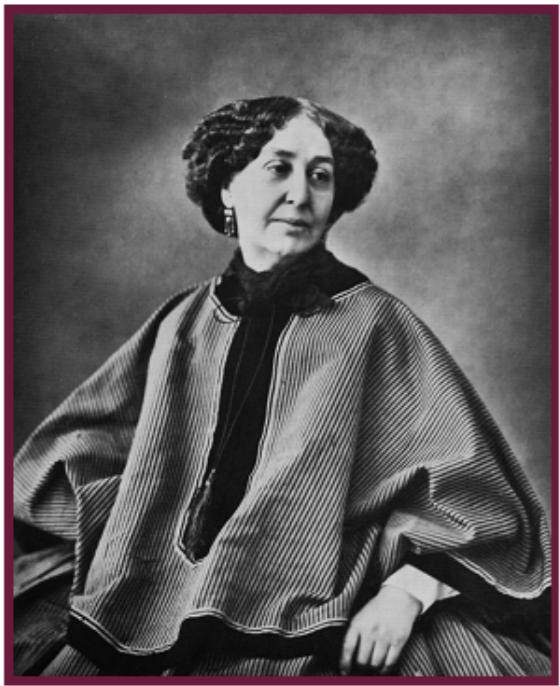
Sept. 29, 2015

Slide Credits:  
Ross T. Whitaker  
Frédo Durand

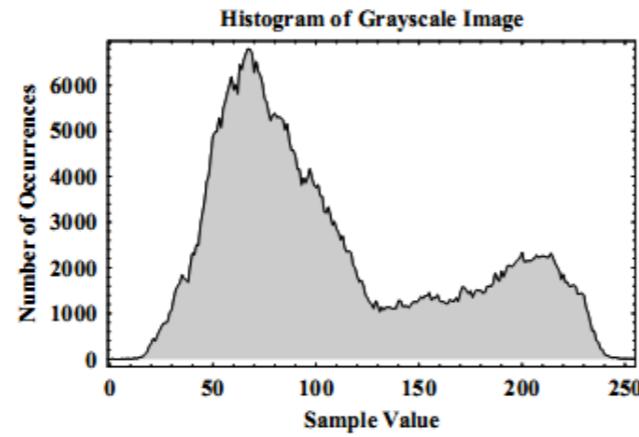
# Agenda

- PA02 Results
- Q02 Due on Thurs.
- PA03 Questions?

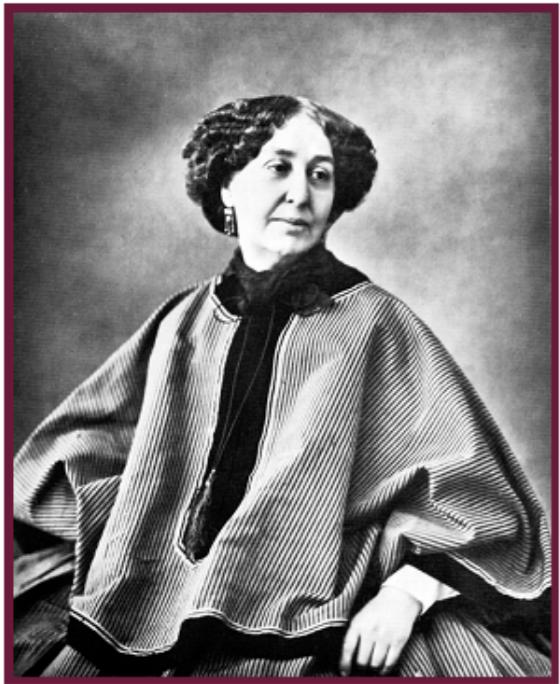
# Last Time



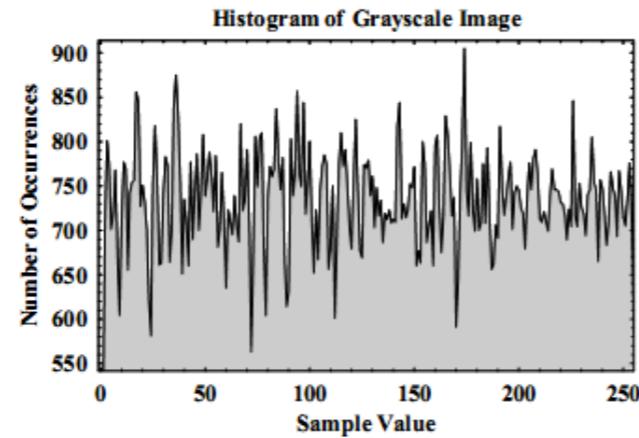
(a) Underexposed image.



(b) Histogram of the underexposed image.



(c) Histogram-equalized image.

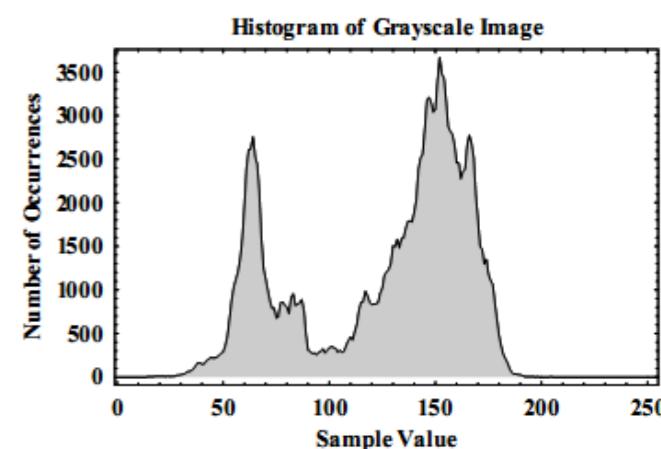


(d) Histogram of the equalized image.

Figure 5.10. Histogram equalization of an underexposed image.



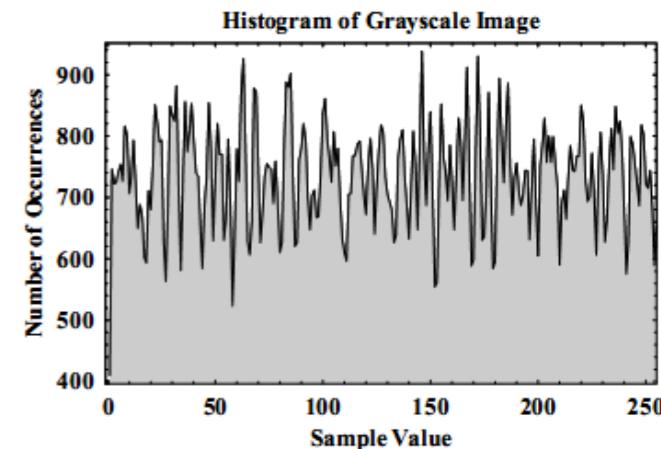
(a) Low-contrast image.



(b) Histogram of the low-contrast image.



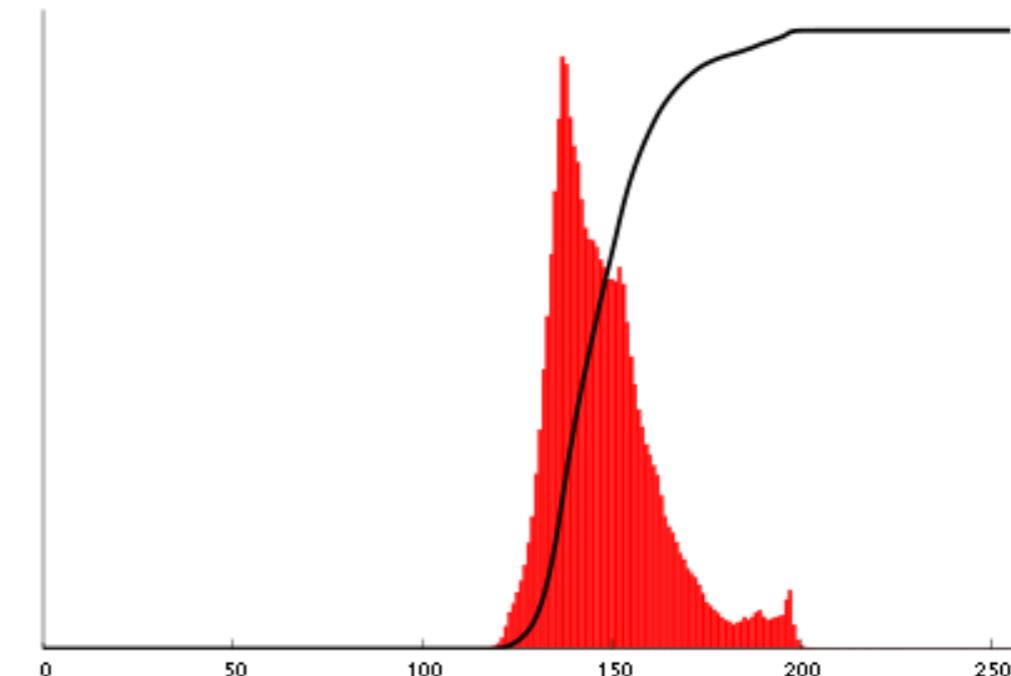
(c) Histogram equalized image.



(d) Histogram of the equalized image.

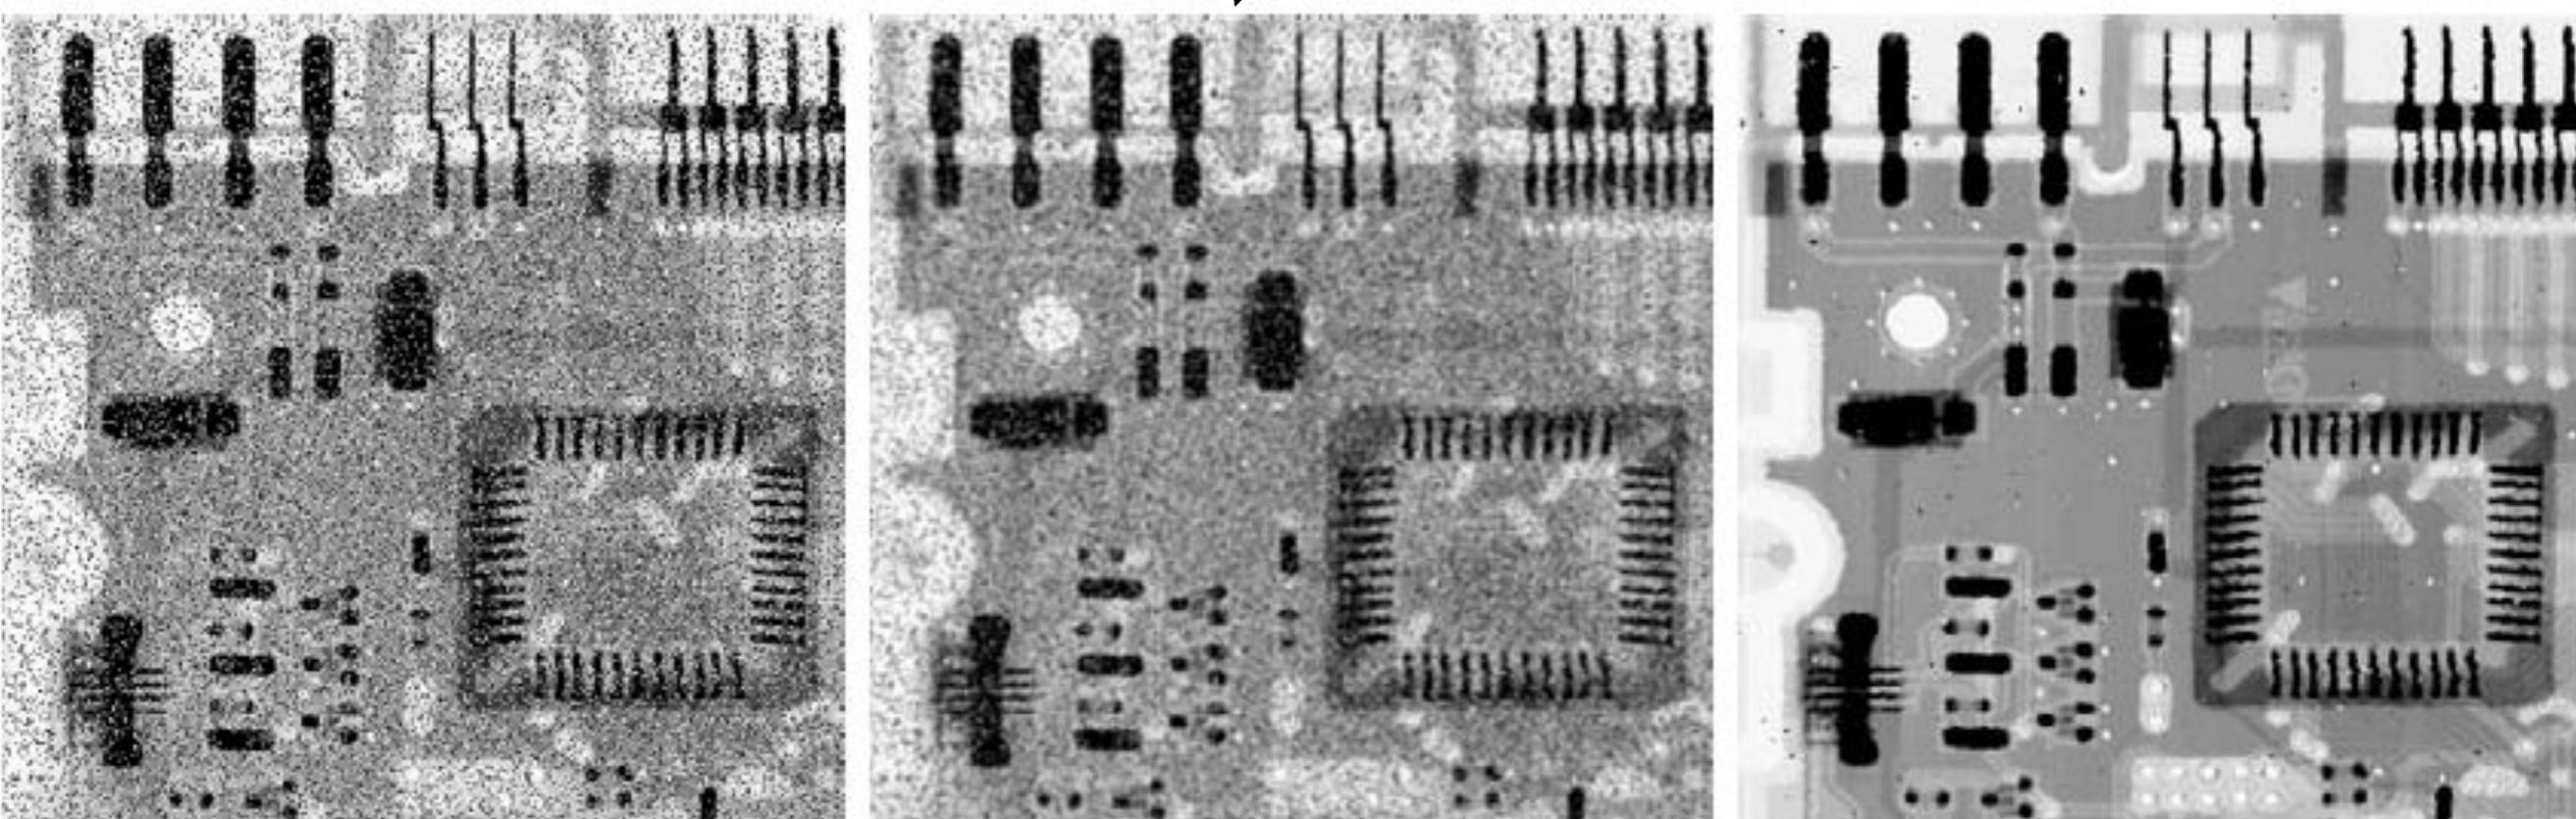
Figure 5.11. Histogram equalization of a low-contrast image.

# The Black Line is the CDF



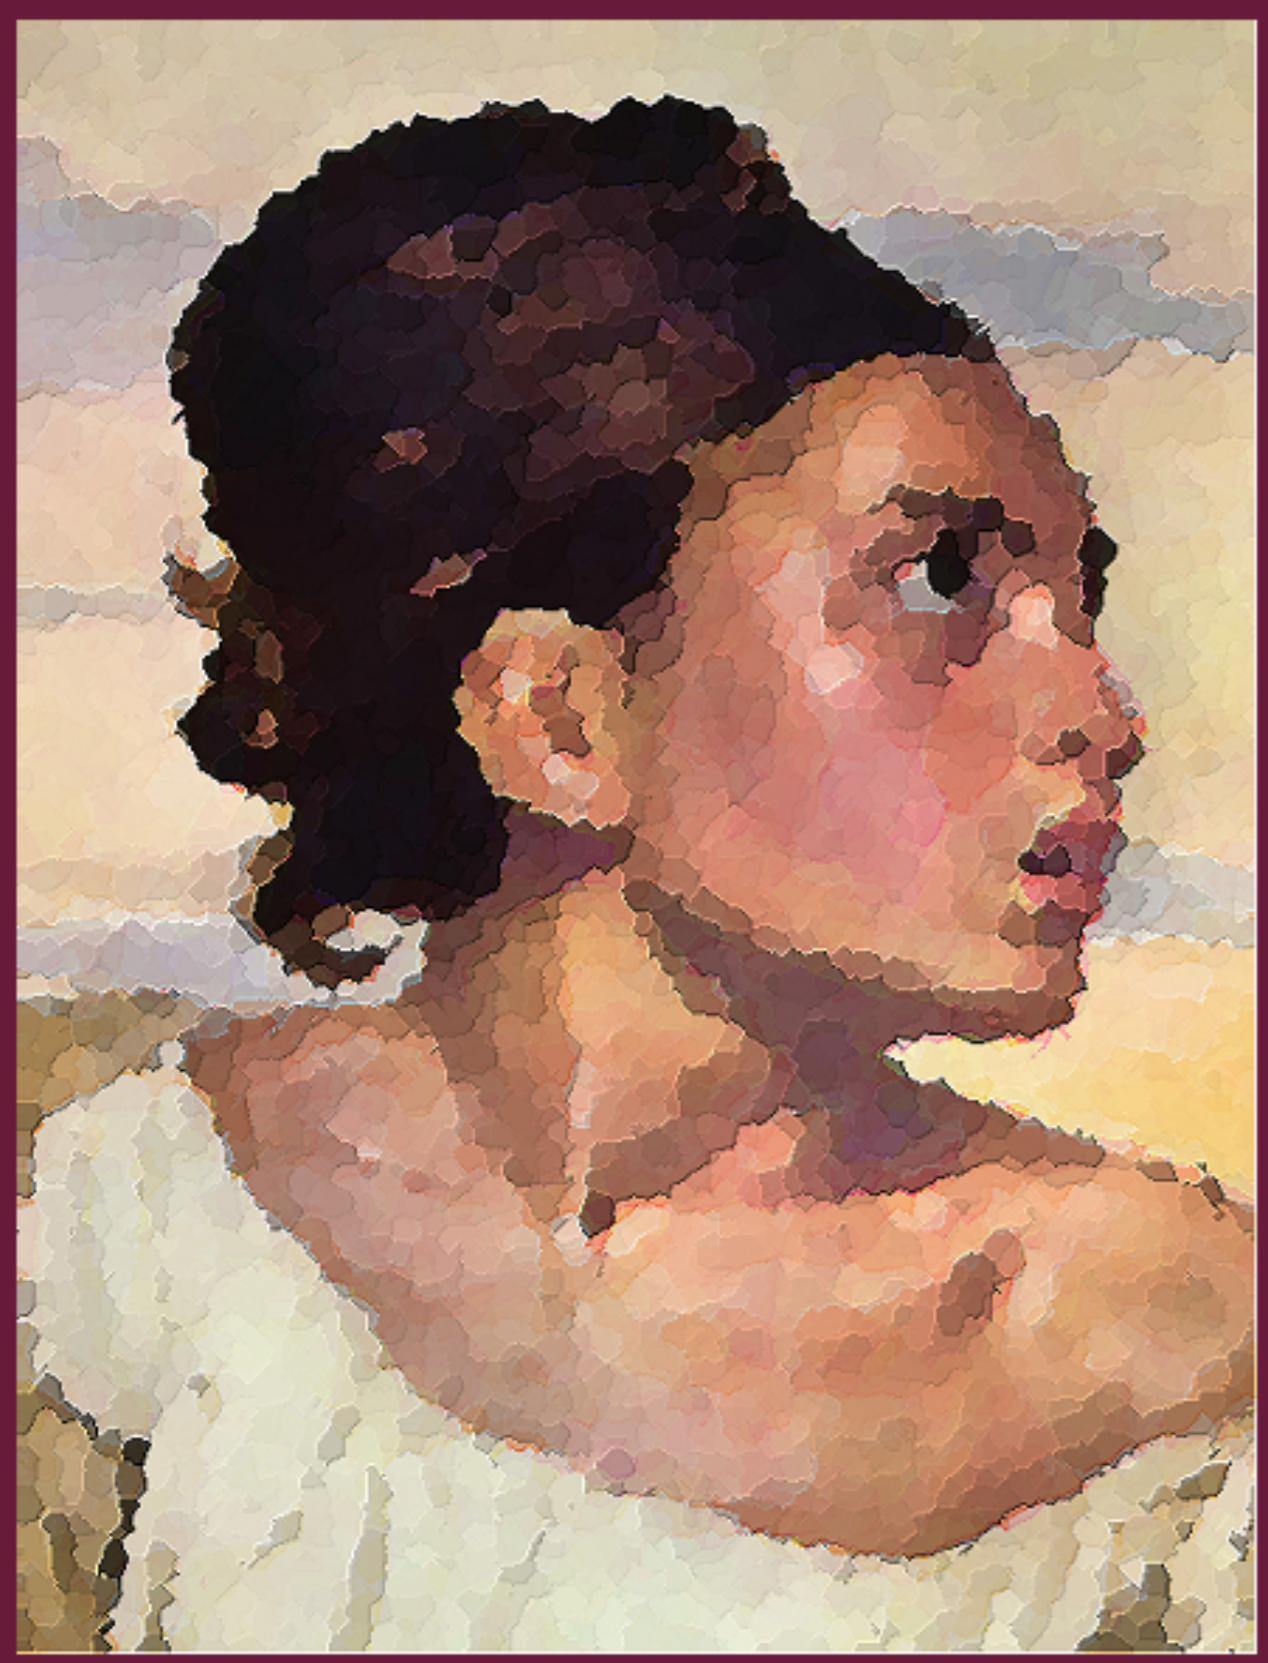
# Median vs. Mean Filtering

Output =  $\text{mean}(N_i)$



a b c

**FIGURE 3.35** (a) X-ray image of circuit board corrupted by salt-and-pepper noise. (b) Noise reduction with a  $3 \times 3$  averaging mask. (c) Noise reduction with a  $3 \times 3$  median filter. (Original image courtesy of Mr. Joseph E. Pascente, Lixi, Inc.)



(a) Most common sample.



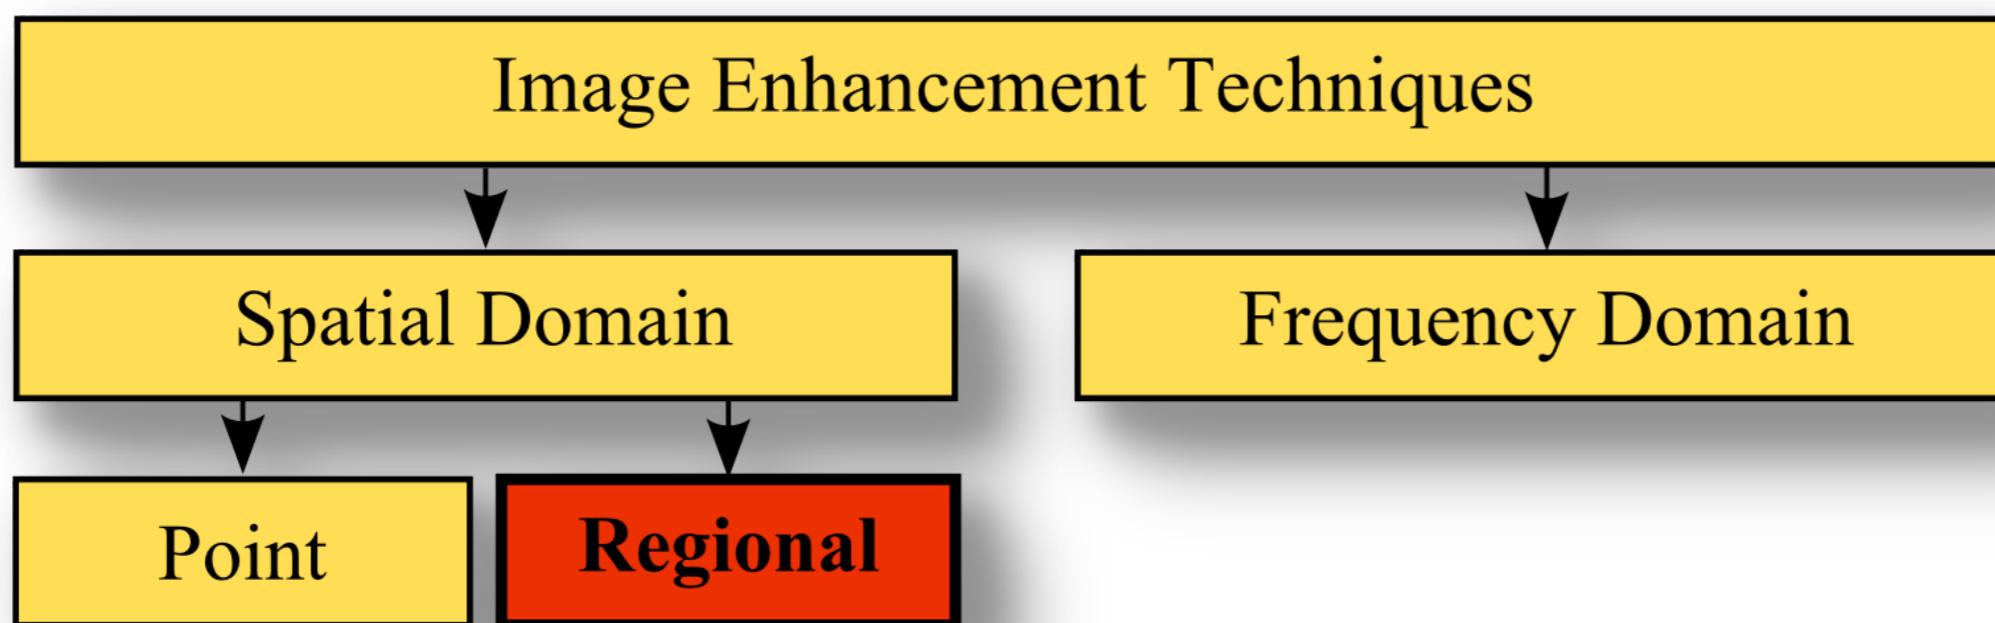
(b) Range filtering.

Figure 6.26. Other statistically-based filters.

# Convolution

# Overview of Region Processing

- Under point processing a single input sample is processed to produce a single output sample.
- In regional processing the value of the output sample is dependent on the values of samples within close proximity to the input sample.
- Exactly how the region is selected and how samples within the region affect the output depends upon the desired effect.
- **Convolution** and **Correlation** are two important regional operations



# An Example: Mean Filtering

- Mean filters sum all of the pixels in a local neighborhood  $N_i$  and divide by the total number, computing the average pixel.
- Said another way, we replace each pixel as a linear combination of its neighbors (with equal weights!)
$$f(N_i) = 1 / |N_i| \sum C_j, \text{ for pixel } j \text{ in } N_i$$

- Where the  $N_i$  is a square, we call these **box** filters

# Box Filtering



$$\begin{matrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{matrix}$$



$$1/9 * \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$



# Box Filtering



$$1/9 * \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

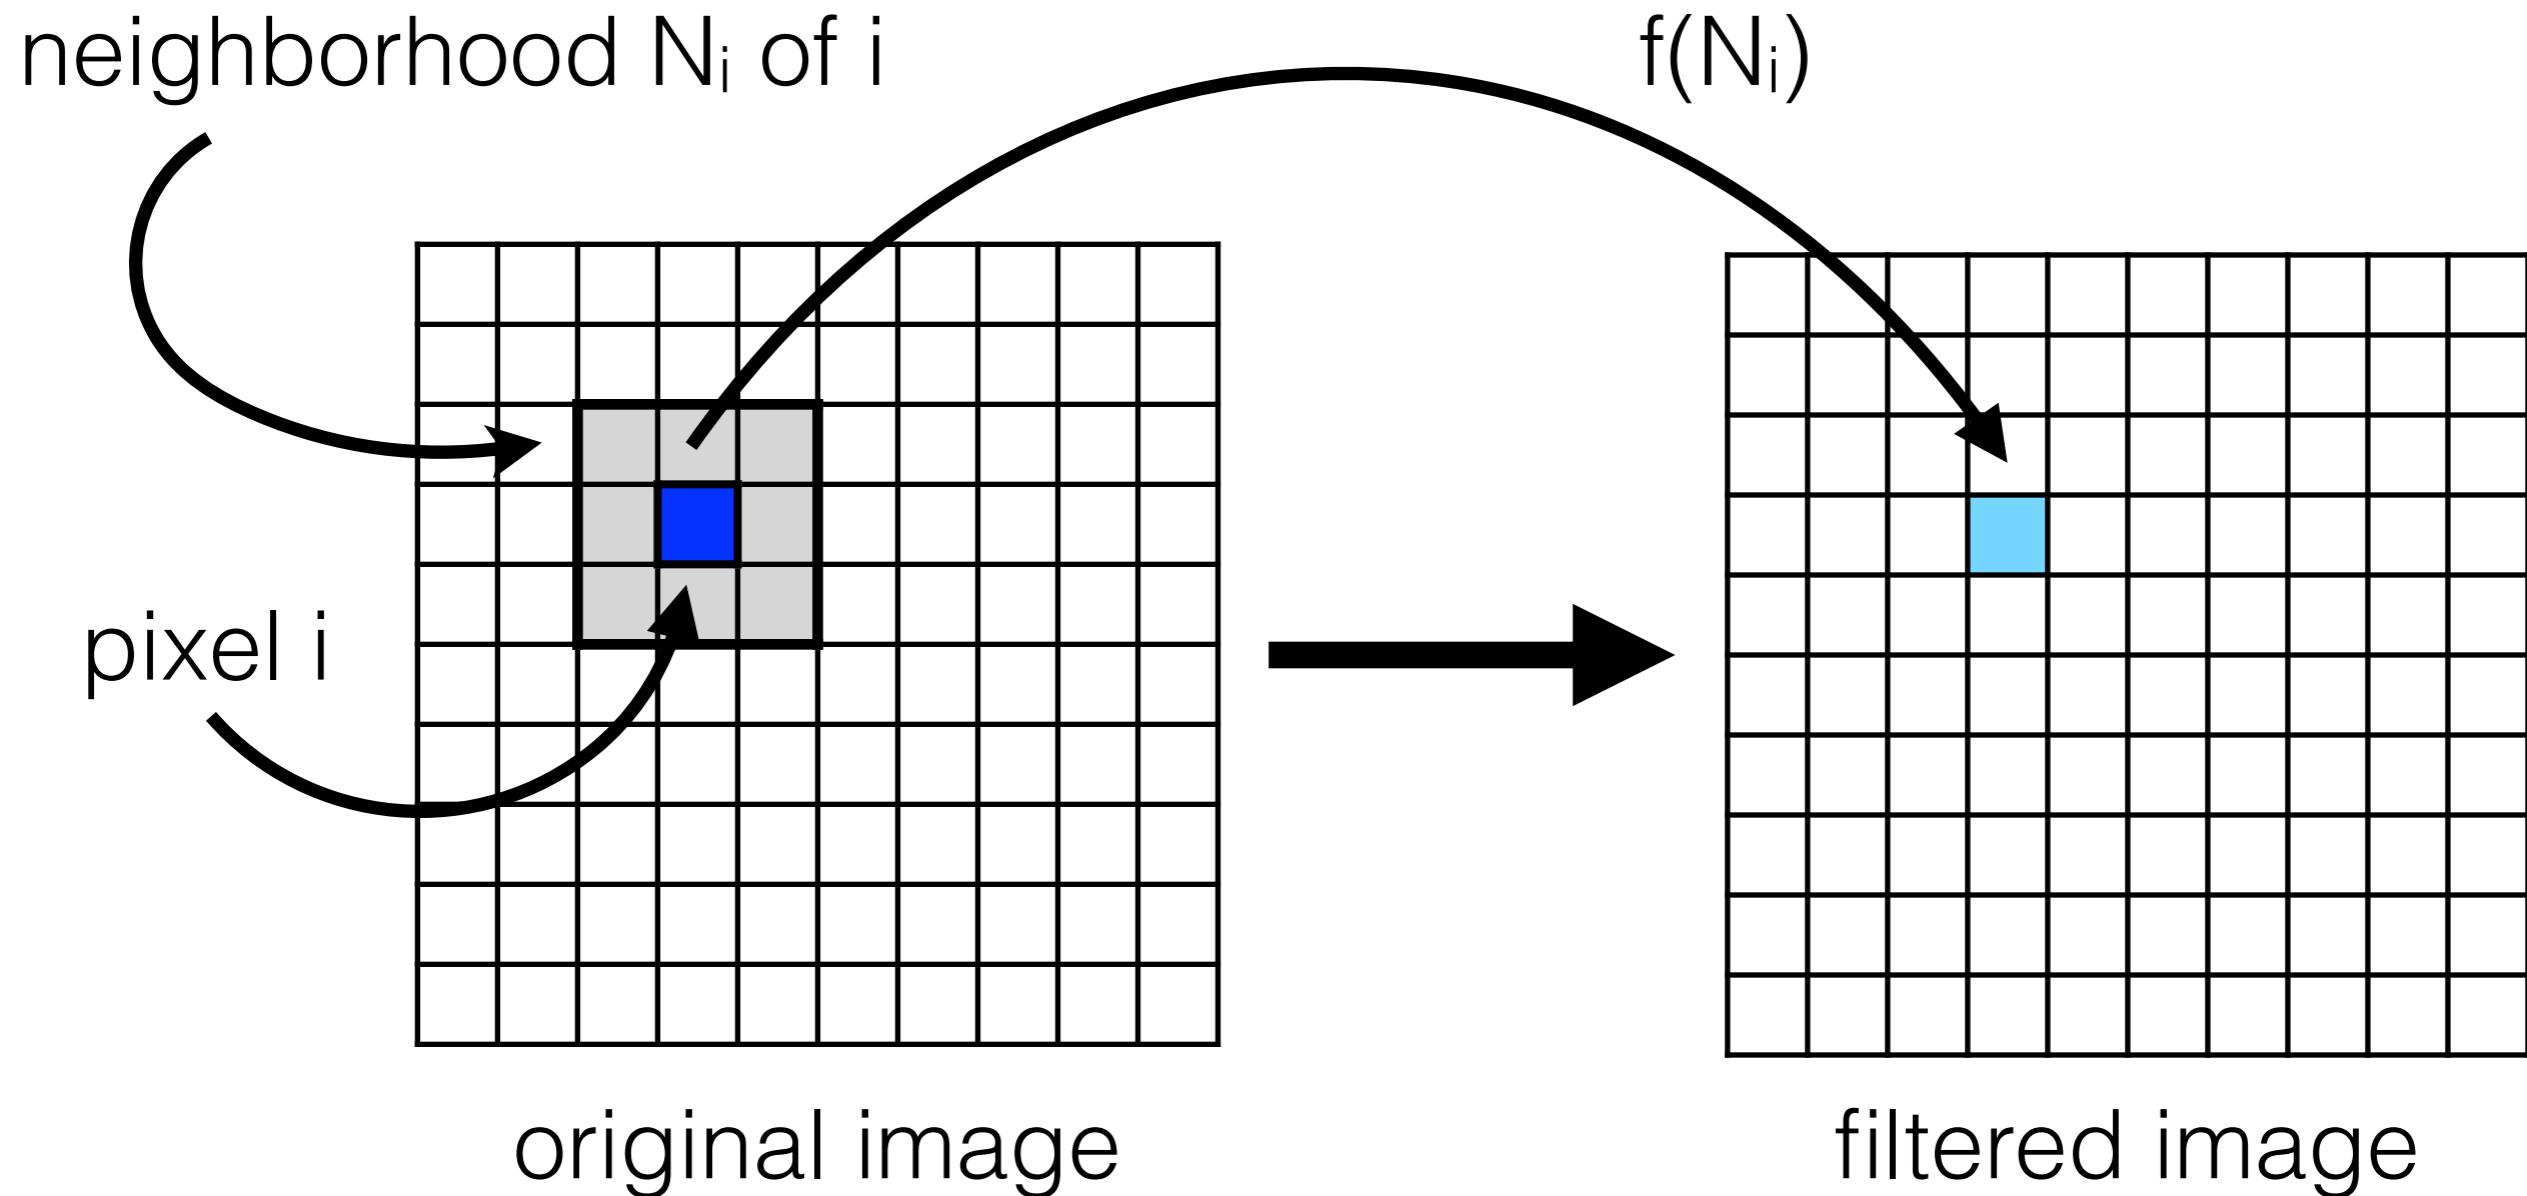


$$1/25 * \begin{matrix} 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \end{matrix}$$



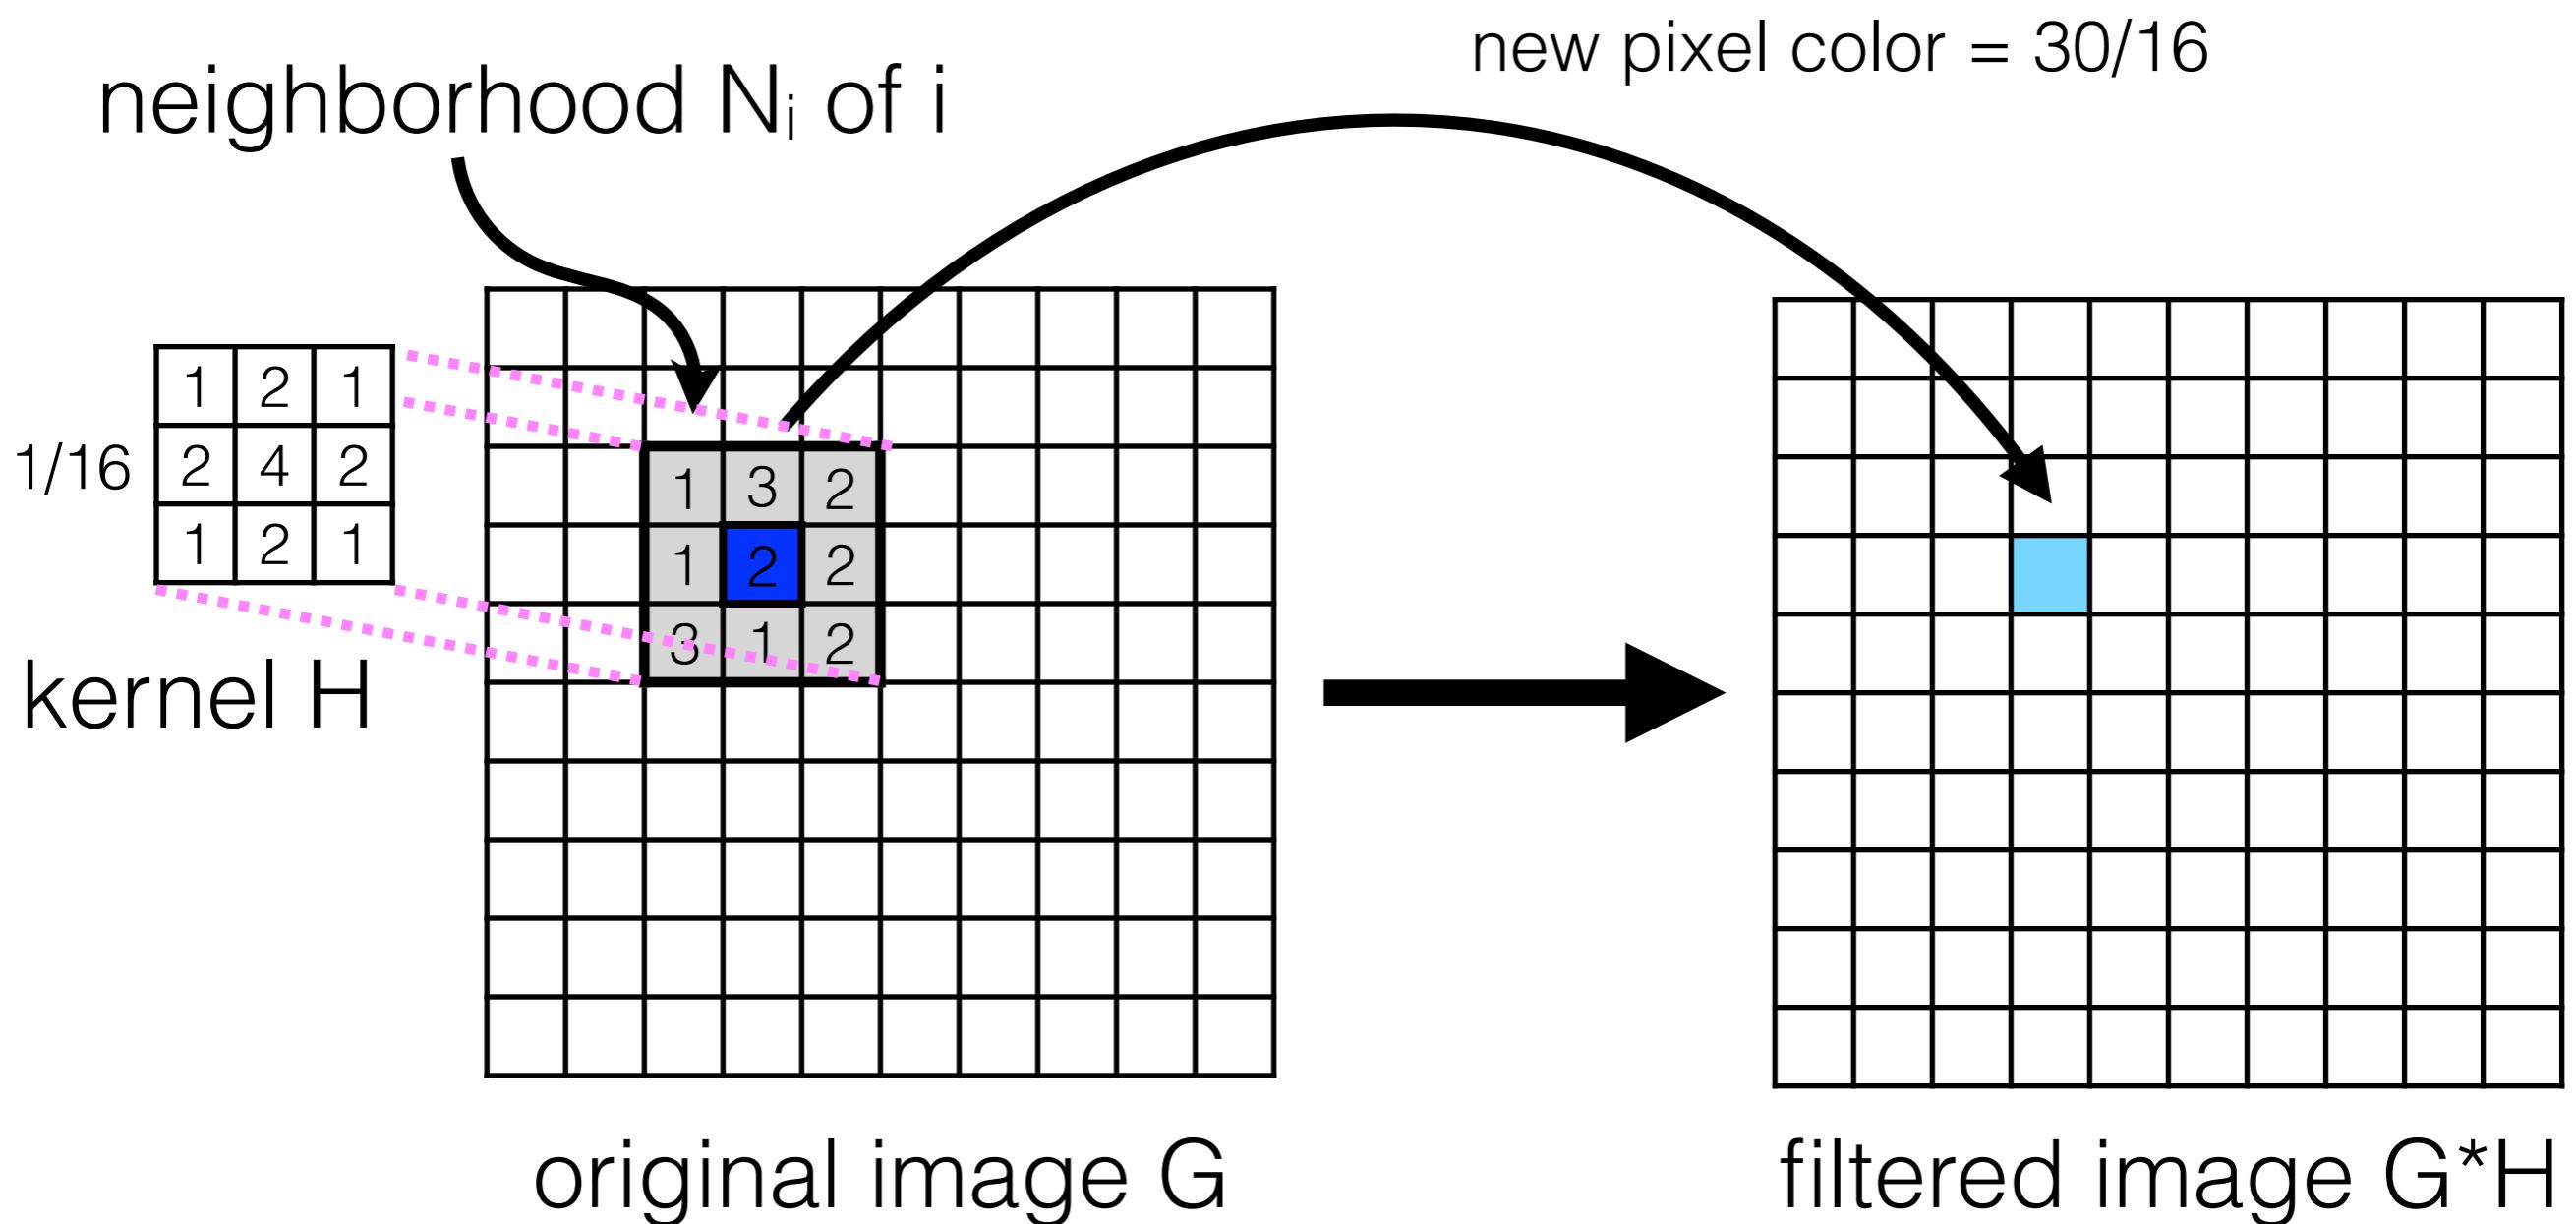
# Recall: Filtering (Schematic)

$$C_{\text{out}} = f(N_{\text{in}})$$



# Convolution

- This process of adding up pixels multiplied by various weights is called **convolution**

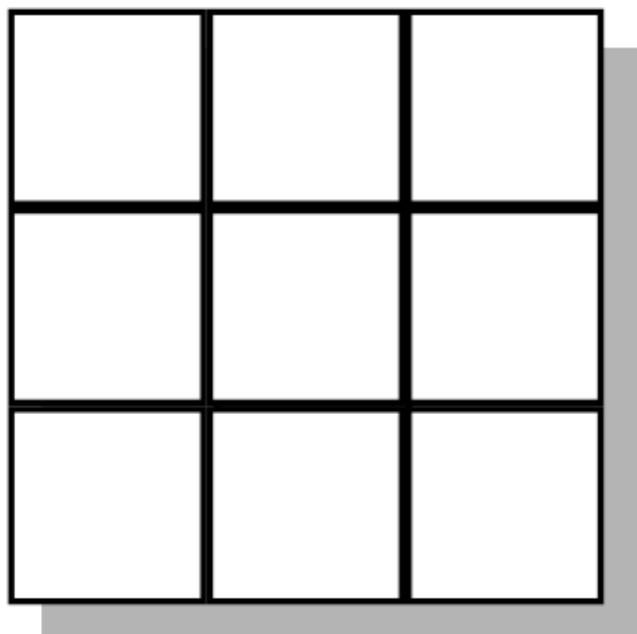


# Kernels

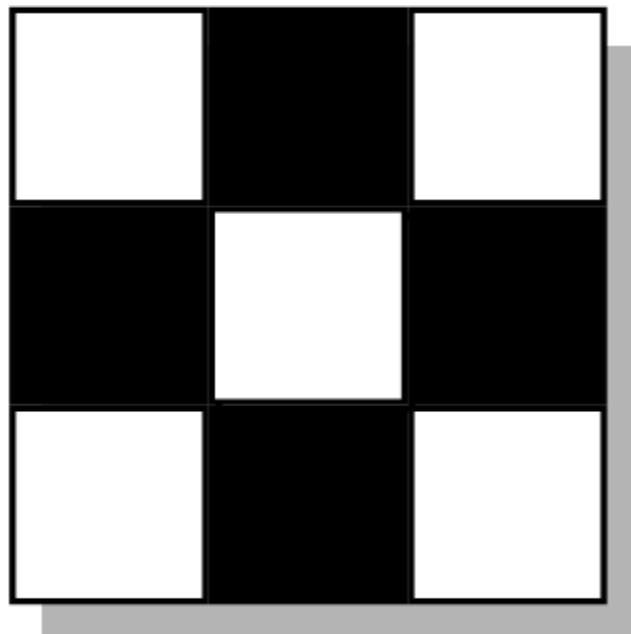
- Convolution employs a rectangular grid of coefficients, known as a **kernel**
- Kernels are like a neighborhood mask, they specify which elements of the image are in the neighborhood and their relative weights.
- A kernel is a set of weights that is applied to corresponding input samples that are summed to produce the output sample.

# Recall: Rank Filtering w/ Masks

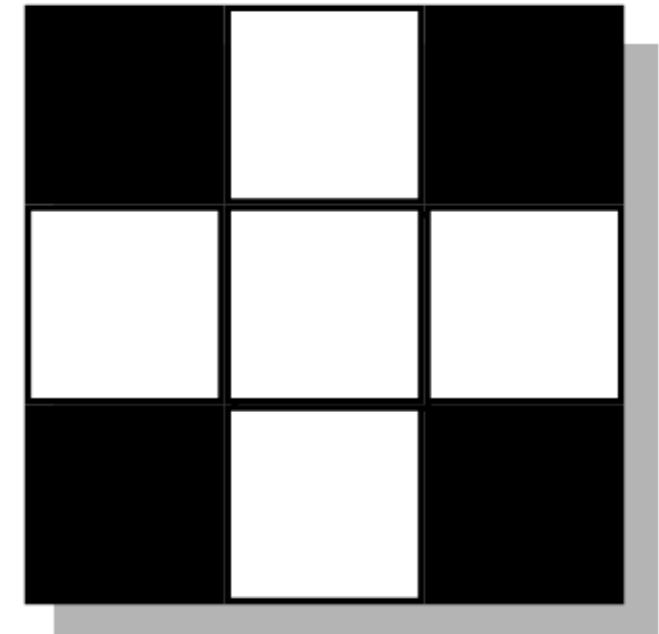
- Rank filtering can also use non-rectangular regions:
  - The most common of such regions are in the shape of either a plus (+) or a cross (x).
- A **mask** is used to specify non-rectangular regions where certain elements in the region are marked as either included (white) or excluded (black) from the region.



(a) Square.



(b) An *x* mask.



(c) A + mask.

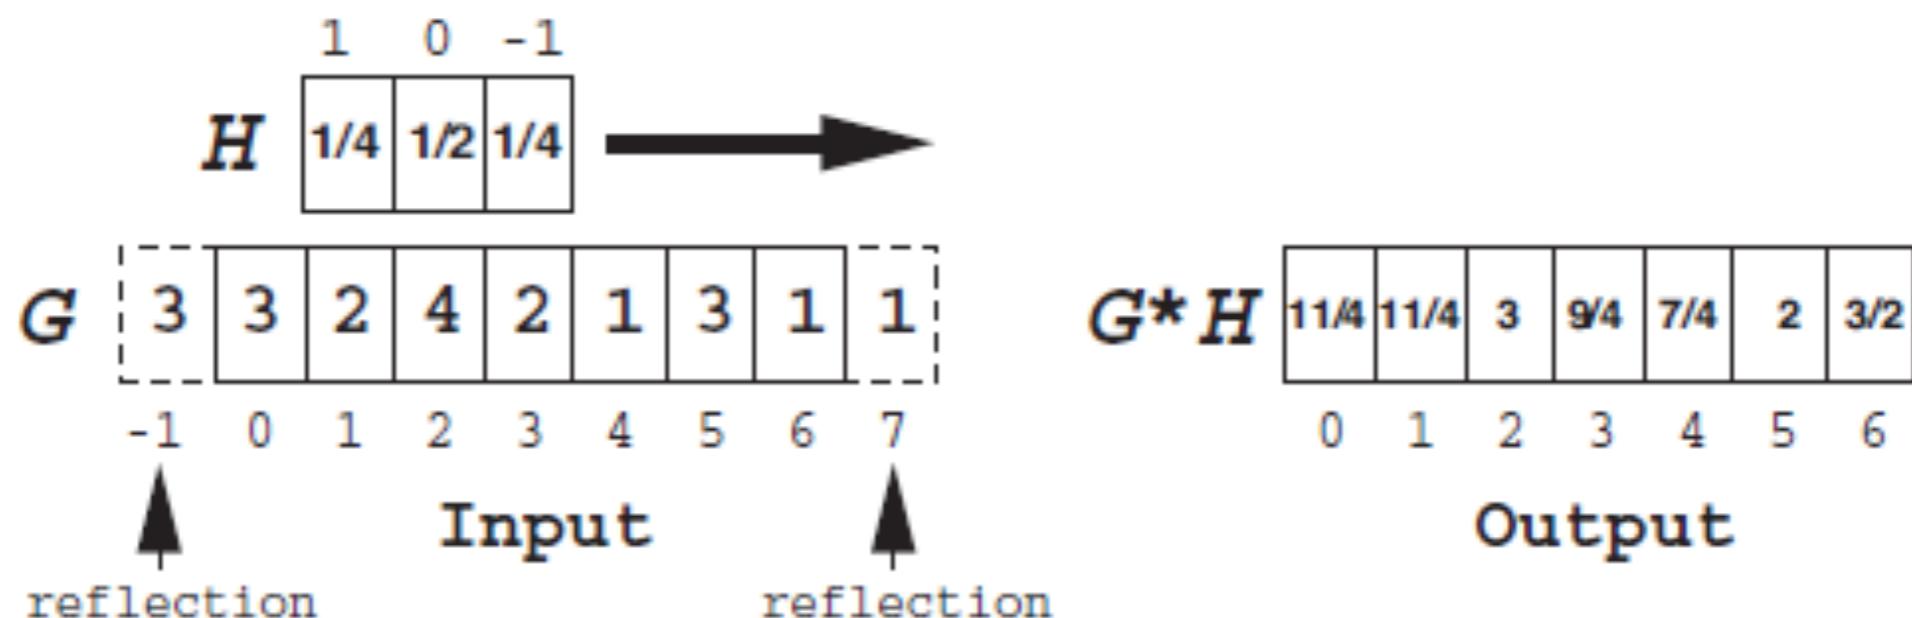
Figure 6.23. Rank filtering masks for a  $3 \times 3$  square,  $3 \times 3$  x, and  $3 \times 3$  +.

# One-dimensional Convolution

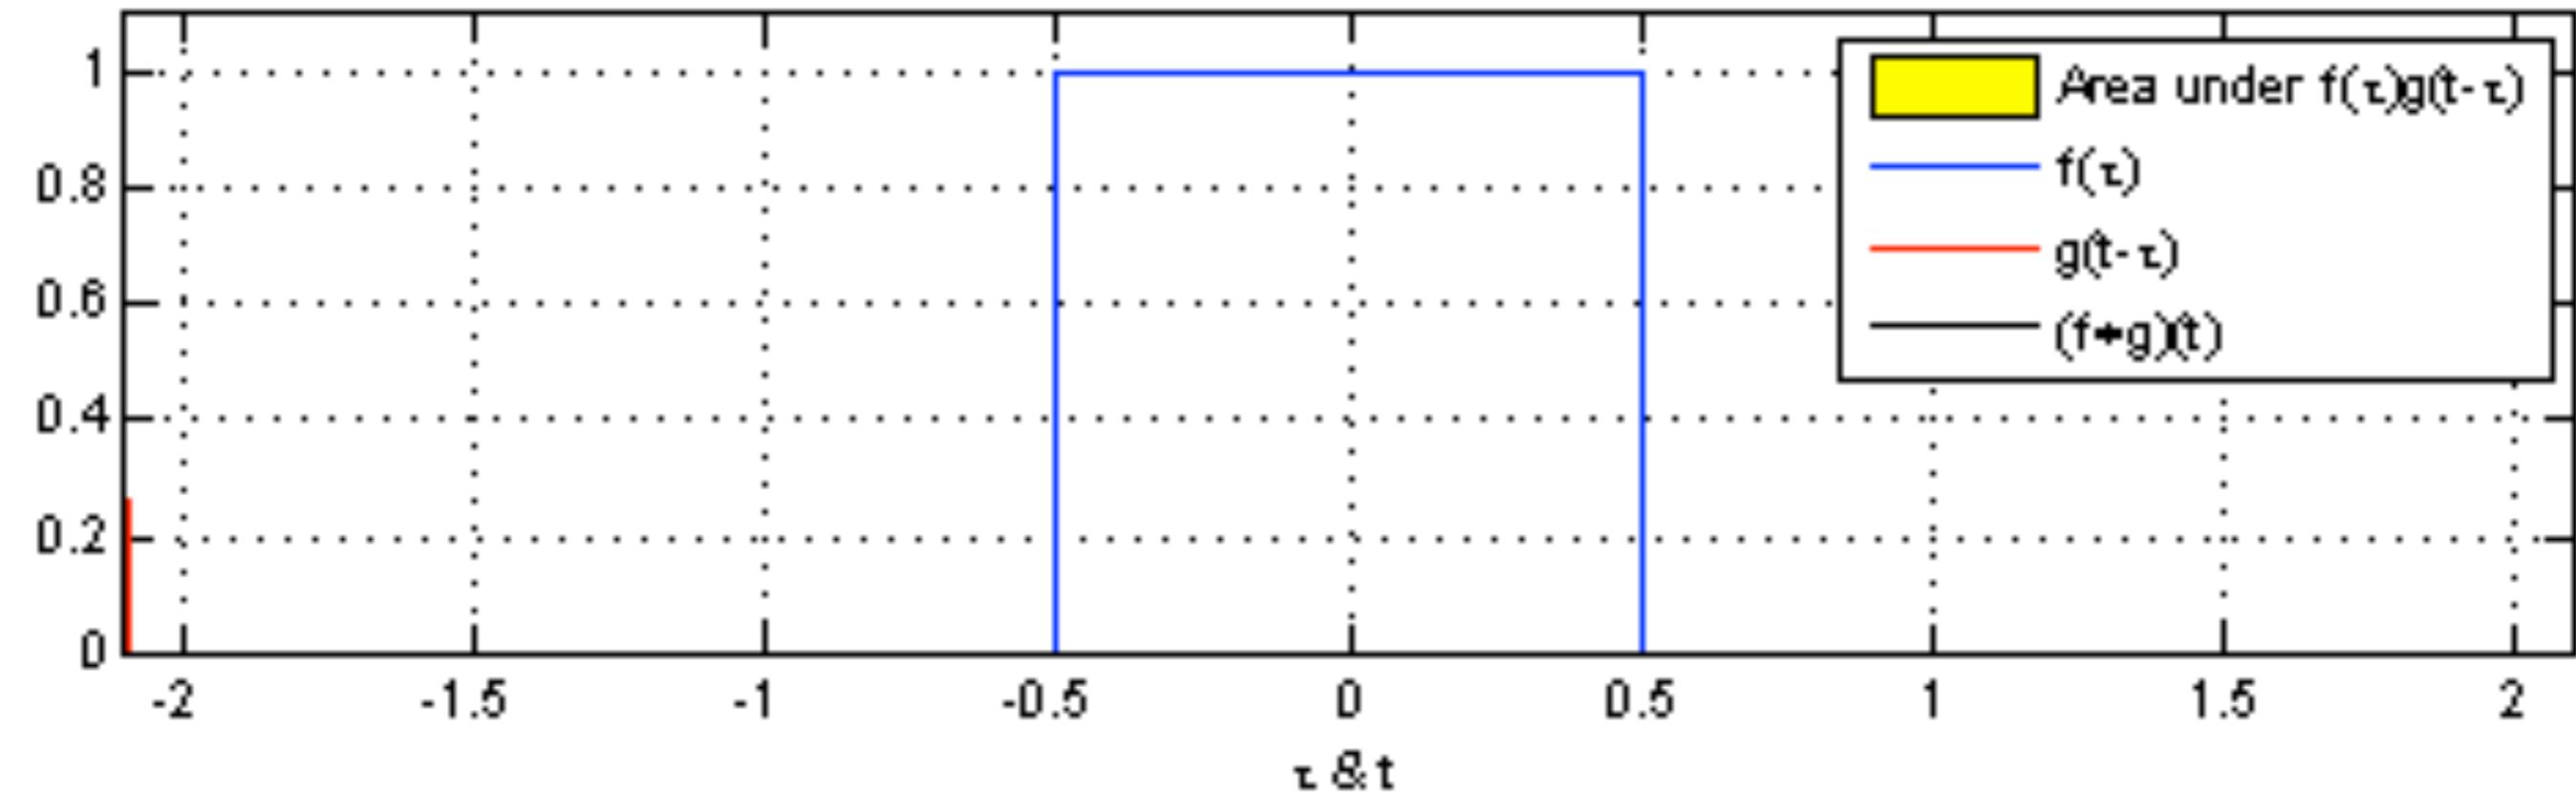
- Can be expressed by the following equation, which takes a filter H and convolves it with G:

$$\hat{G}[i] = (G * H)[i] = \sum_{j=i-n}^{i+n} G[j]H[i-j], \quad i \in [0, N-1]$$

- Equivalent to sliding a window



# Box Filters (Animated)



- The above is continuous, but a discrete version could be imagined as:

$$H = 111$$

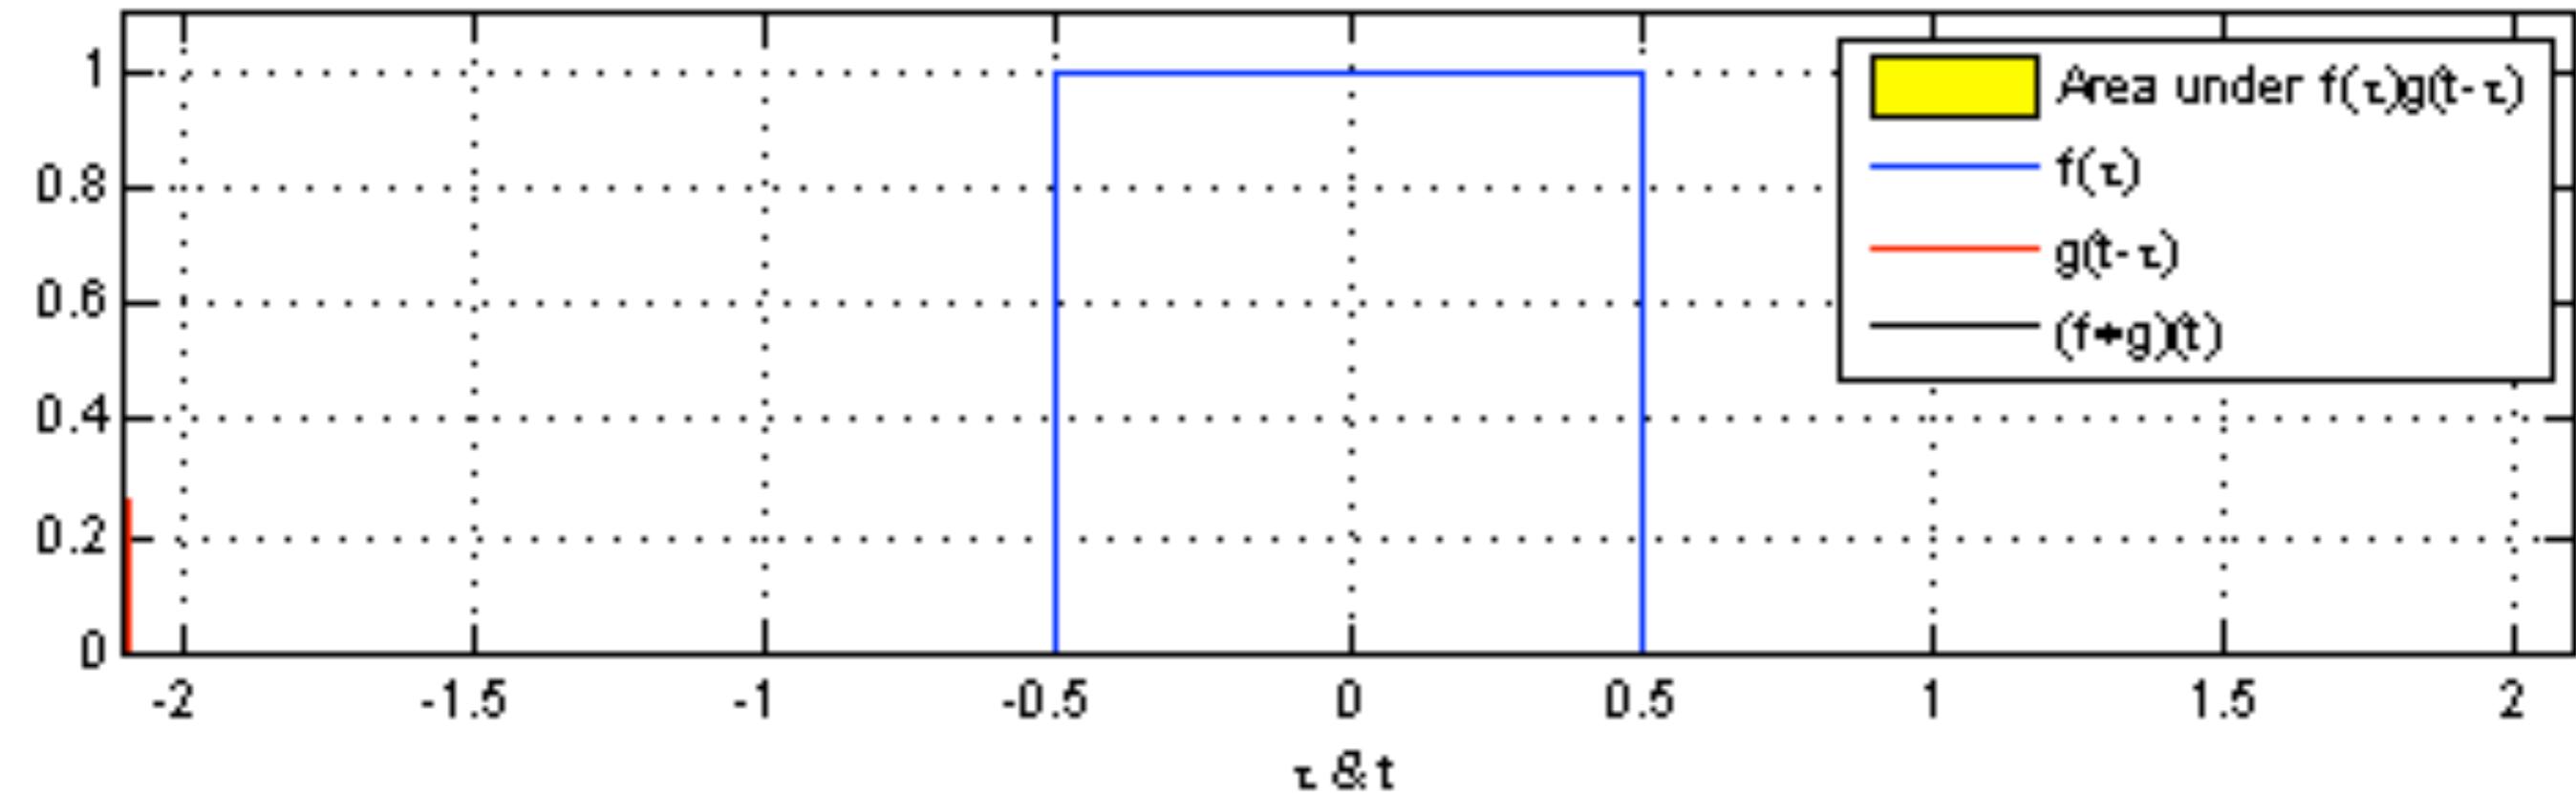
$$G = \dots 0000011100000\dots$$

$$G^*H = \dots 0000012100000\dots$$

<http://en.wikipedia.org/wiki/Convolution>

More examples: <http://math.mit.edu/daimp/ConvFlipDrag.html>

# Box Filters (Animated)



- The above is continuous, but a discrete version could be imagined as:

$$H = 111$$

$$G = \dots 0000011100000\dots$$

$$G^*H = \dots 0000012100000\dots$$

<http://en.wikipedia.org/wiki/Convolution>

More examples: <http://math.mit.edu/daimp/ConvFlipDrag.html>

# 2-Dimensional Version

- Given a  $W \times H$  grayscale image  $I$  and an  $M \times N$  kernel  $K$  such that  $M$  and  $N$  are odd the convolution of  $I$  with  $K$  is given below where  $x$  is in  $[0, W-1]$  and  $y$  is in  $[0, H-1]$ .

$$(I \otimes K)(x, y) = I'(x, y) = \sum_{j=-\lfloor M/2 \rfloor}^{\lfloor M/2 \rfloor} \sum_{k=-\lfloor N/2 \rfloor}^{\lfloor N/2 \rfloor} K(j, k) \cdot I(x - j, y - k). \quad (6.1)$$

- The  $(x-j)$  and  $(y-k)$  terms can be understood as reflections of the kernel about the central vertical and horizontal axes.
- The kernel weights are multiplied by the corresponding image samples and then summed together.

# A note on indexing

- Convolution **reflects** the filter to preserve orientation.
  - Correlation does **not** have this reflection.
    - But we often use them interchangeably since most kernels are symmetric.

Convolution **reflects**  
and **shifts** the kernel

Given kernel  $H = \begin{matrix} & 1 & 2 & 3 \\ 4 & & 5 & 6 \\ & 7 & 8 & 9 \end{matrix}$

# An Illustration

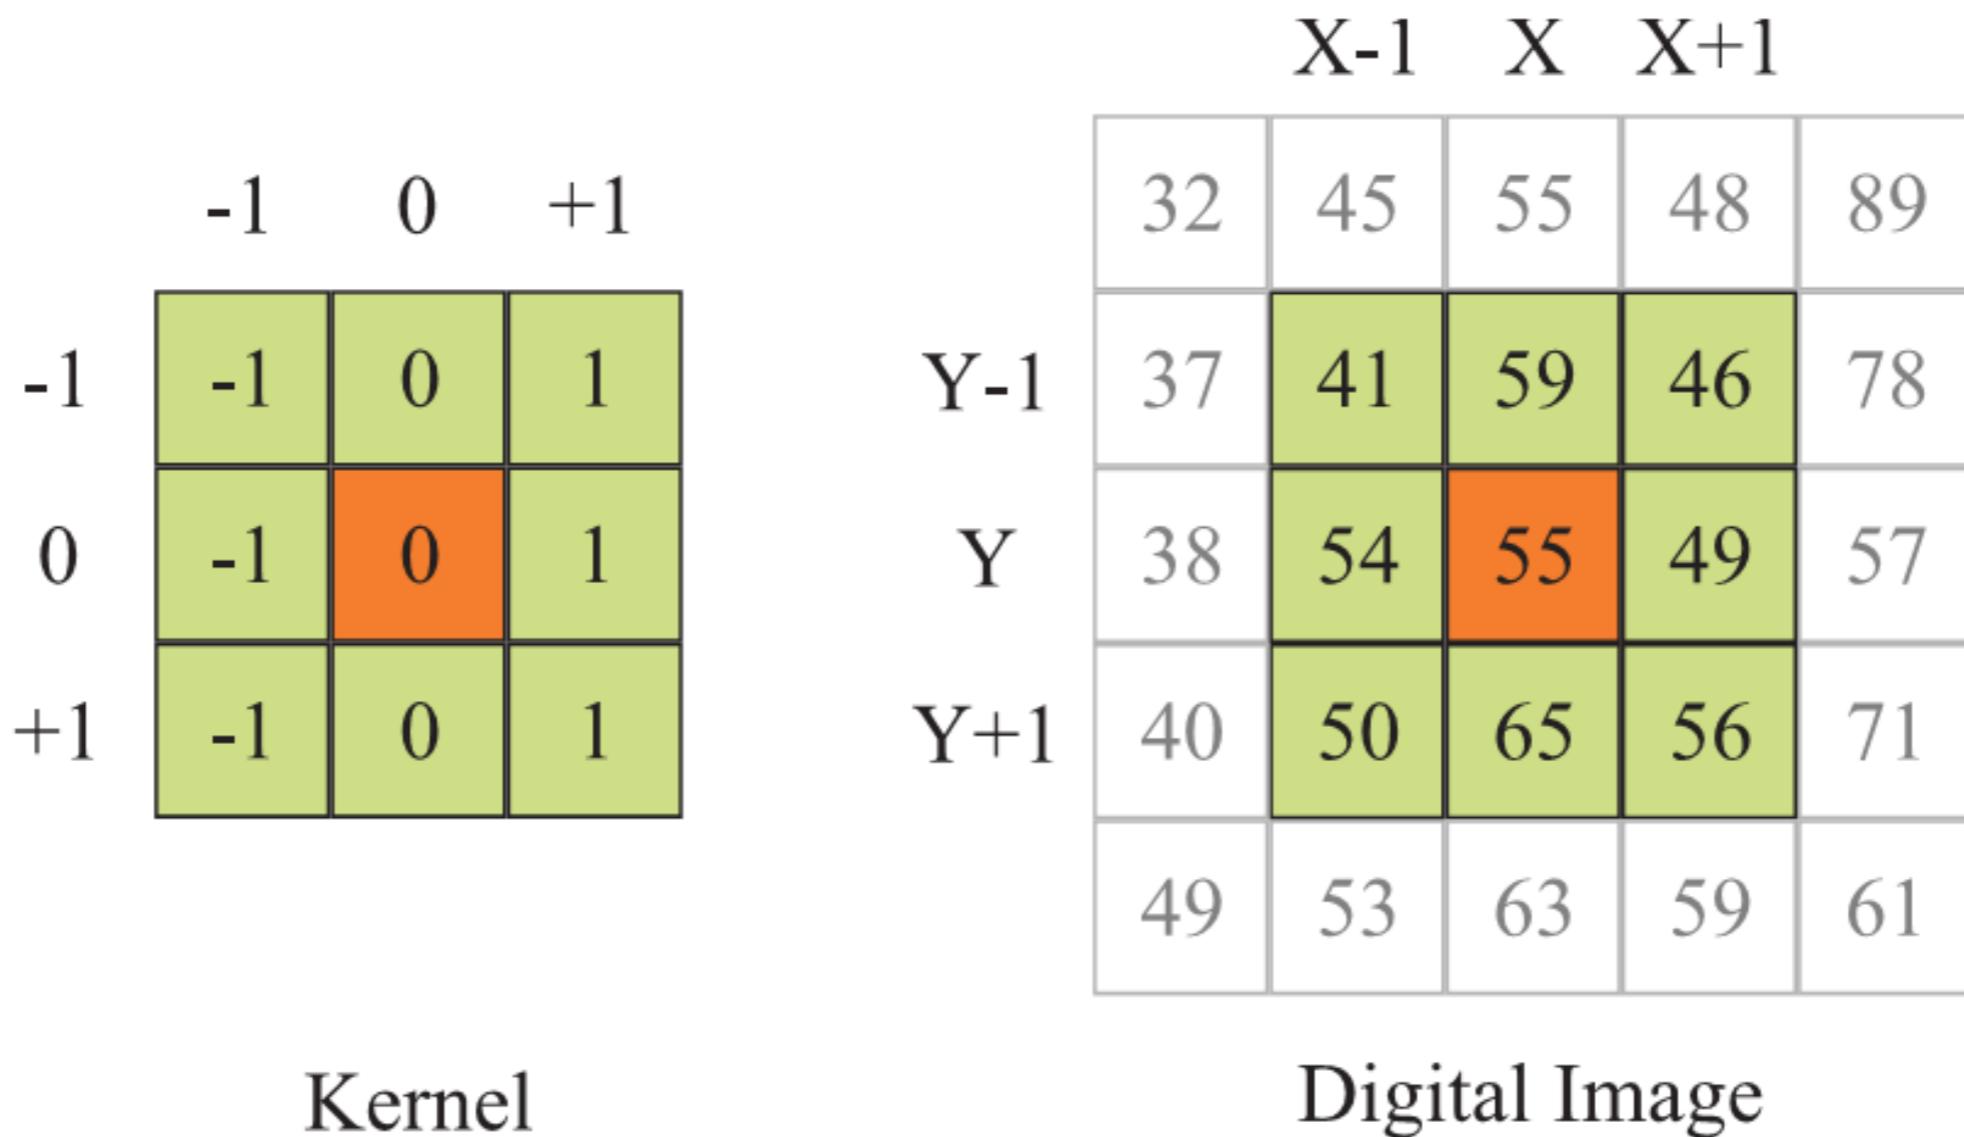


Figure 6.2. A  $3 \times 3$  kernel is centered over sample  $I(x, y)$ .

# Illustration

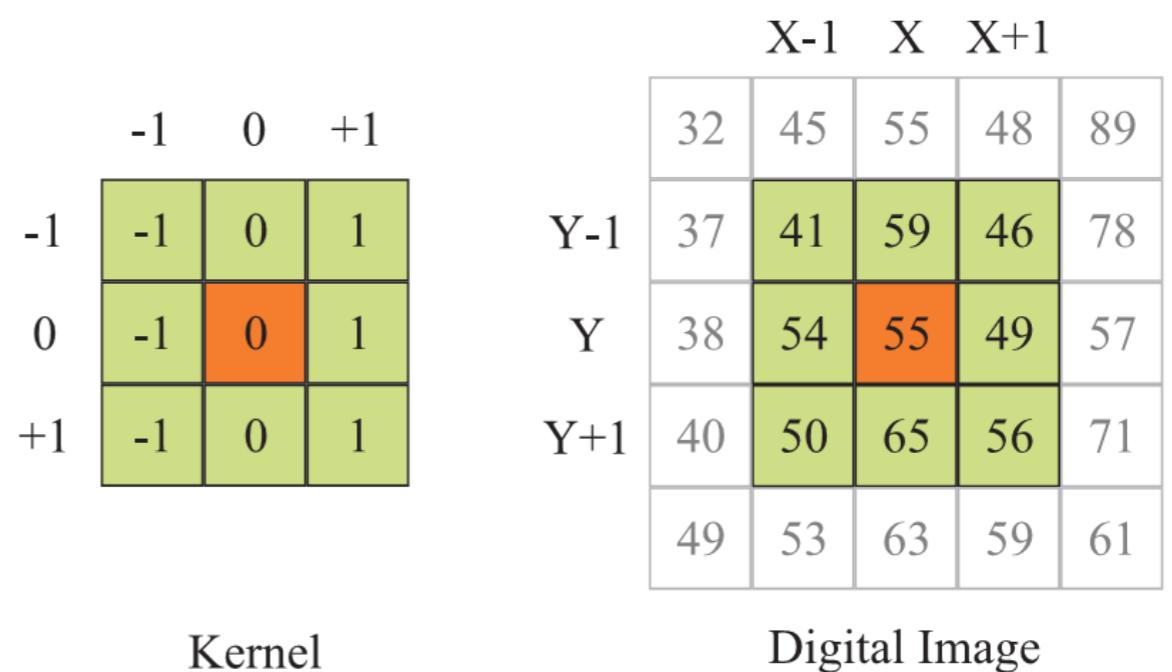
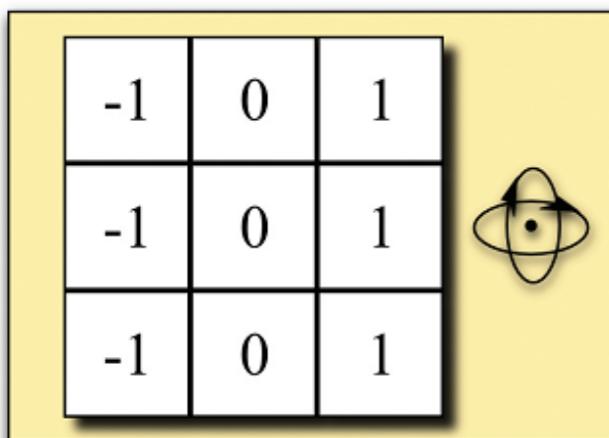
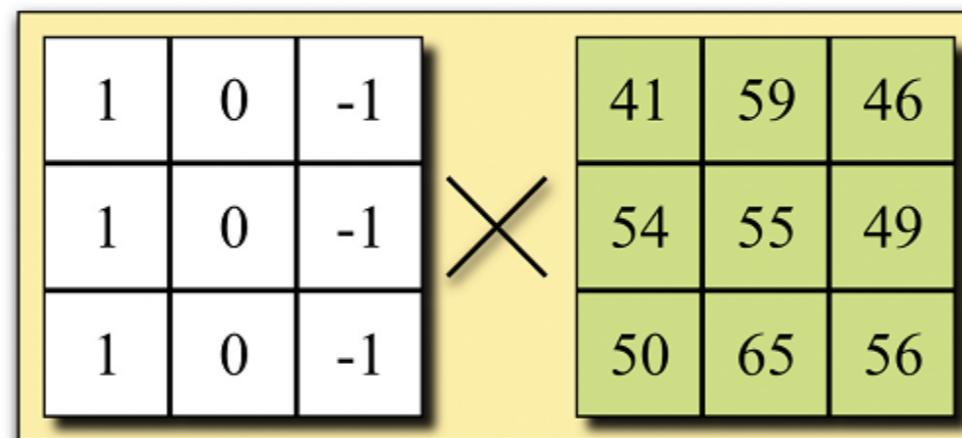


Figure 6.2. A  $3 \times 3$  kernel is centered over sample  $I(x, y)$ .

Reflect Kernel



Term by Term Multiplication



Sum the Products

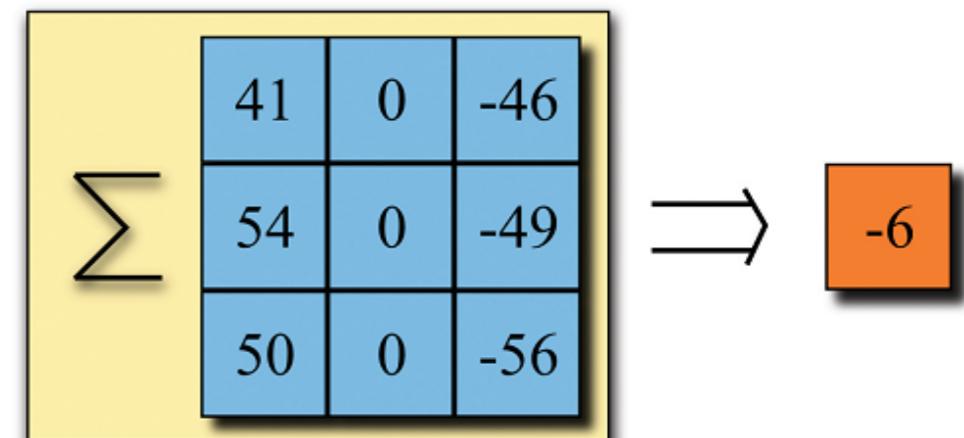
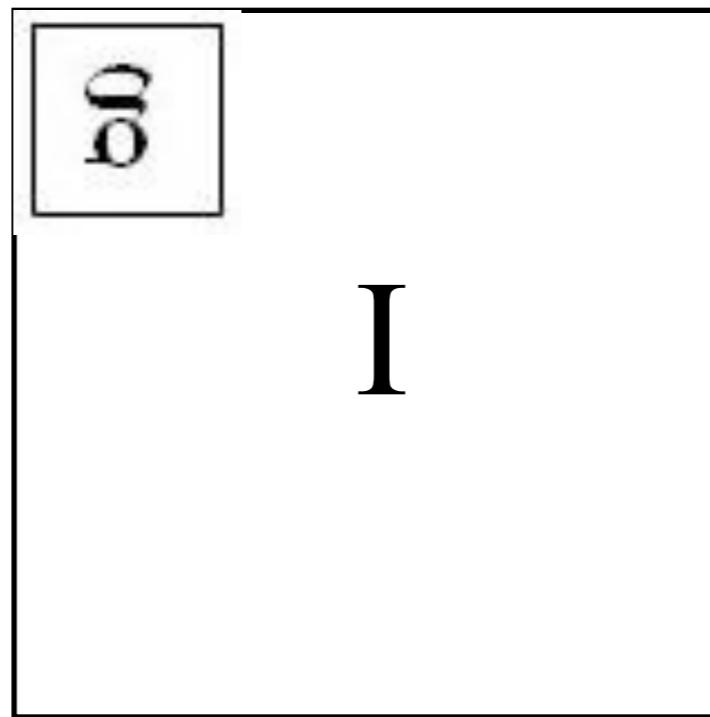


Figure 6.3. Convolution steps.

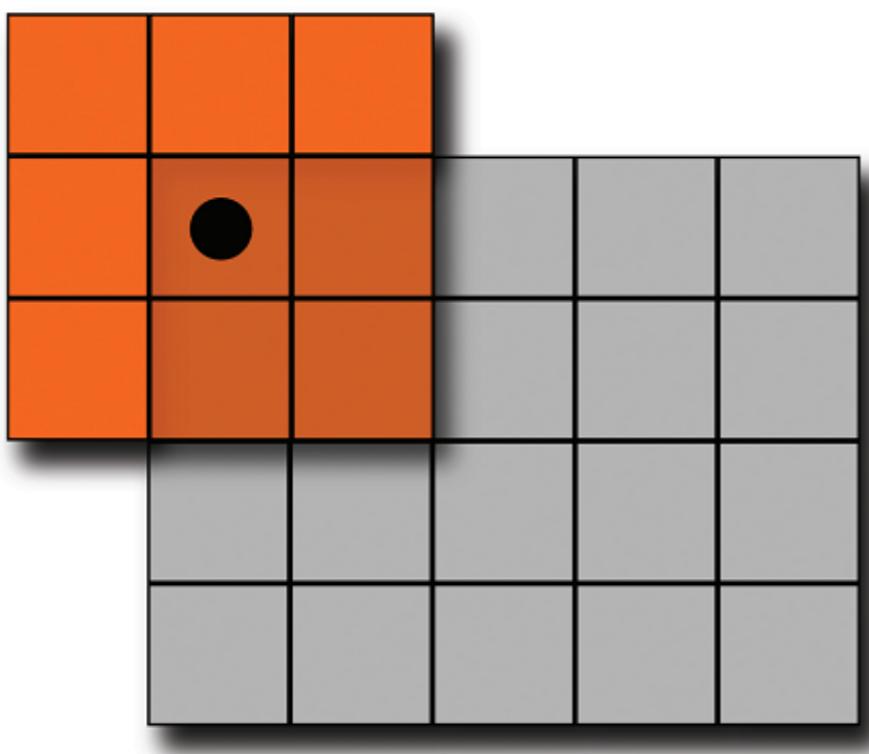
# More Formally

$$I \otimes g(x) = \int_{x'} I(x') g(x - x')$$

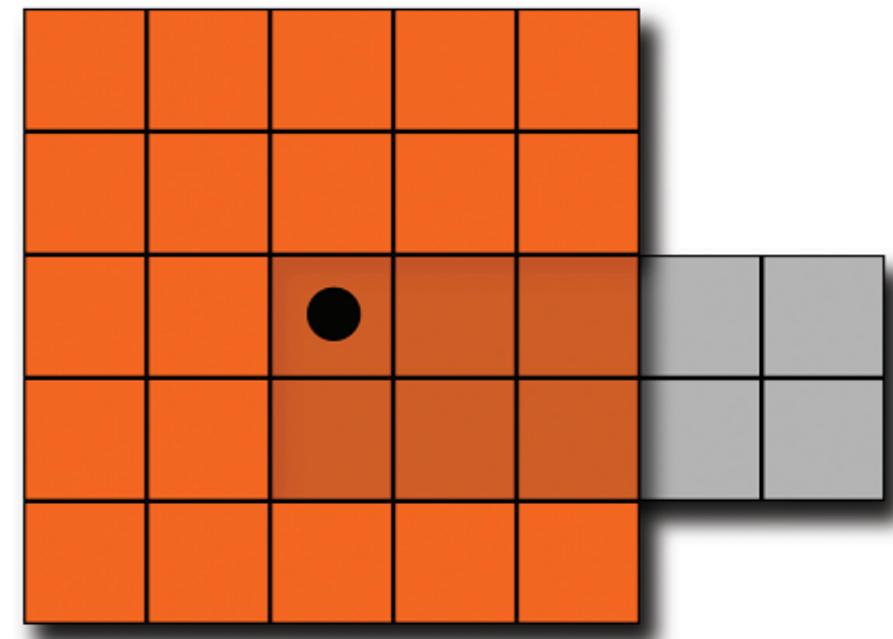


# Reminder: Boundaries

- Consider how to handle boundaries. What should be done if the kernel falls off of the boundary of the source image as shown in the illustrations below?



(a) Kernel at  $I(0, 0)$ .

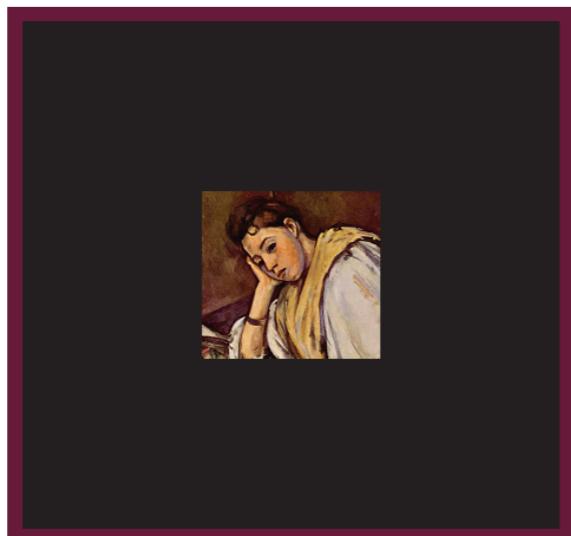


(b) Kernel larger than the source.

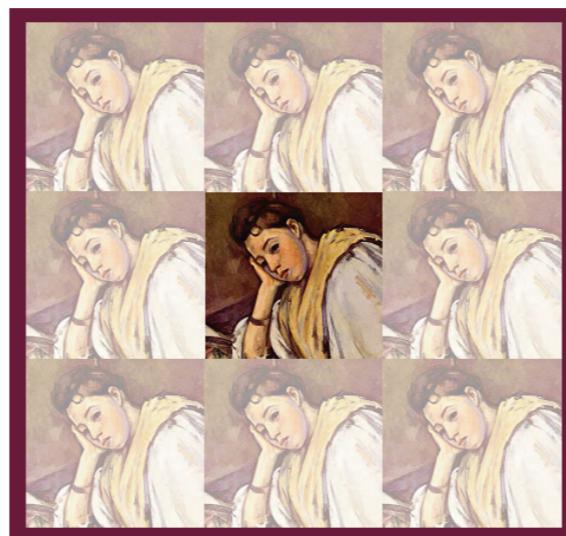
Figure 6.4. Illustration of the edge handling problem.

# Boundary Padding

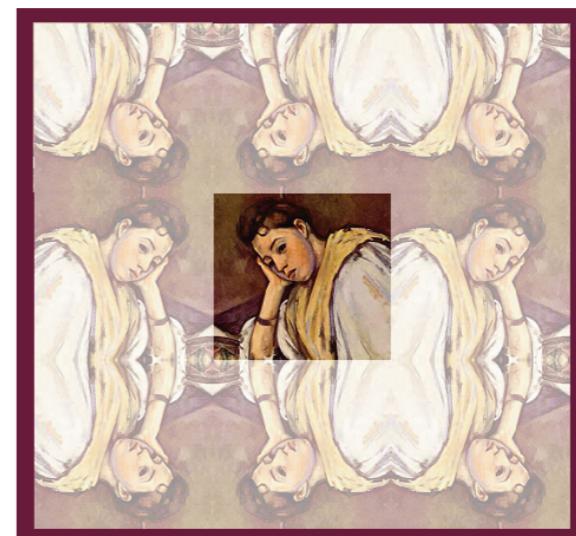
- Either we redefine convolution to:
  - Produce zero when the kernel falls off of the boundary. If the kernel extends beyond the source image when centered on a sample  $I(x, y)$  then the output sample is set to zero.
  - Produce  $I(x, y)$  when the kernel falls off the boundary. If the kernel extends beyond the source image when centered on a sample  $I(x, y)$  then the output sample is defined as  $I(x, y)$ .
- Or else we use boundary padding:



(a)



(b)



(c)

Figure 6.5. (a) Zero padding, (b) circular indexing, and (c) reflected indexing.

# Kernel Rescaling

- Important: rescale the kernel by making the coefficients sum to 1.
  - The effect of the kernel is unchanged
  - Rescaling the kernel is equivalent to rescaling the convolved image.

1	1	1
1	1	1
1	1	1

(a) Non-normalized.

$$\frac{1}{9} \bullet \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array} = \begin{array}{|c|c|c|} \hline .111 & .111 & .111 \\ \hline .111 & .111 & .111 \\ \hline .111 & .111 & .111 \\ \hline \end{array}$$

(b) Normalization.

.111	.111	.111
.111	.111	.111
.111	.111	.111

(c) Normalized.

Figure 6.6. Normalizing a kernel.

# Smoothing Filters

# Smoothing Spatial Filters

- Any weighted filter with positive values will smooth in some way, examples:

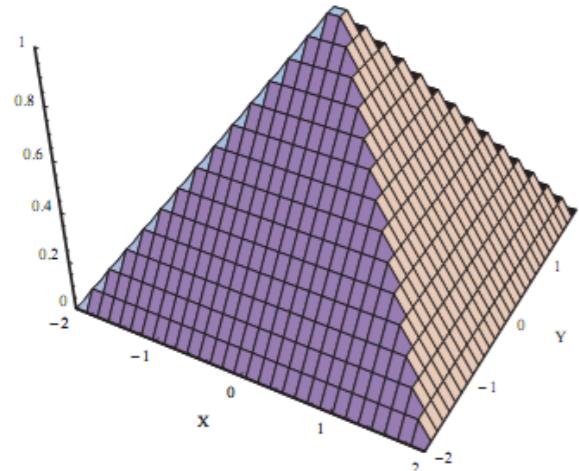
$$\frac{1}{9} \times \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$
$$\frac{1}{16} \times \begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 2 & 4 & 2 \\ \hline 1 & 2 & 1 \\ \hline \end{array}$$

- Normally, we use integers in the filter, and then divide by the sum (computationally more efficient)
- These are also called **blurring** or **low-pass** filters

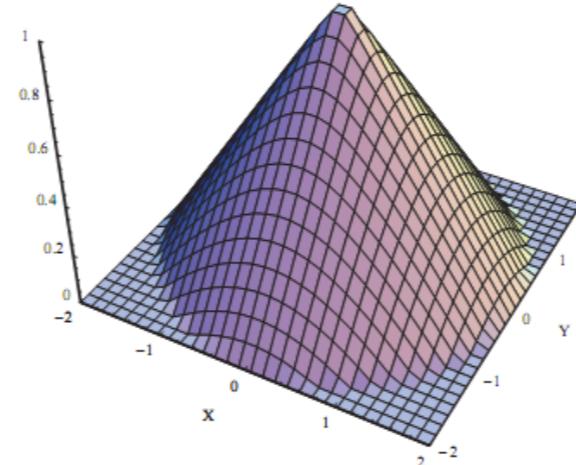
# Smoothing Kernels

$$f(x, y) = -\alpha \cdot \max(|x|, |y|)$$

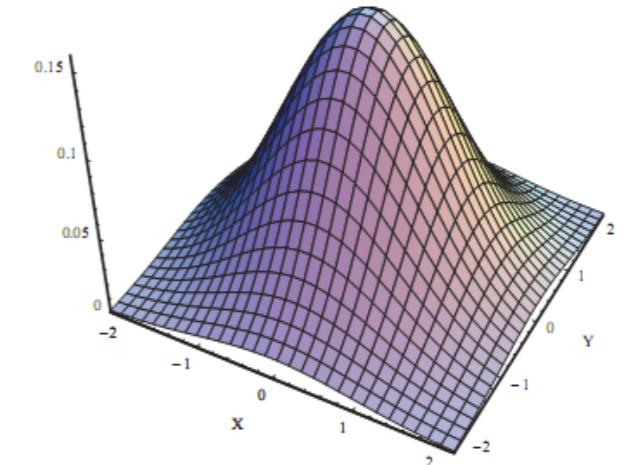
$$G(x, y) = \frac{1}{2\pi\sigma^2} e^{\frac{-(x^2+y^2)}{2\sigma^2}}$$



(a) Pyramid.



(b) Cone.



(c) Gaussian.

$$f(x, y) = -\alpha \cdot \sqrt{x^2 + y^2}$$

1	2	3	2	1
2	4	6	4	2
3	6	9	6	3
2	4	6	4	2
1	2	3	2	1

(a) Pyramid.

0	0	1	0	0
0	2	2	2	0
1	2	5	2	1
0	2	2	2	0
0	0	1	0	0

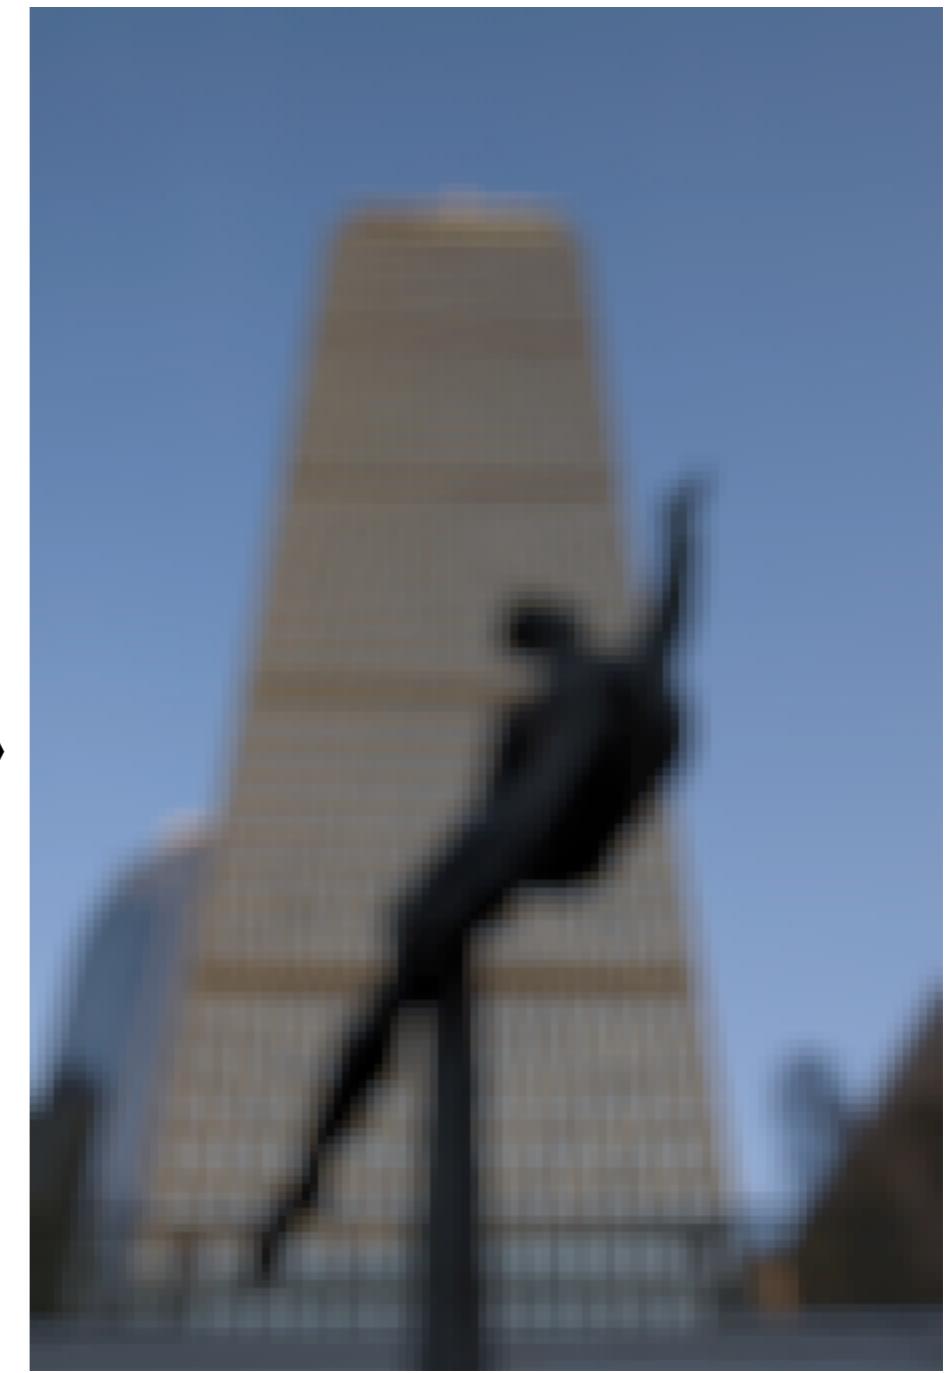
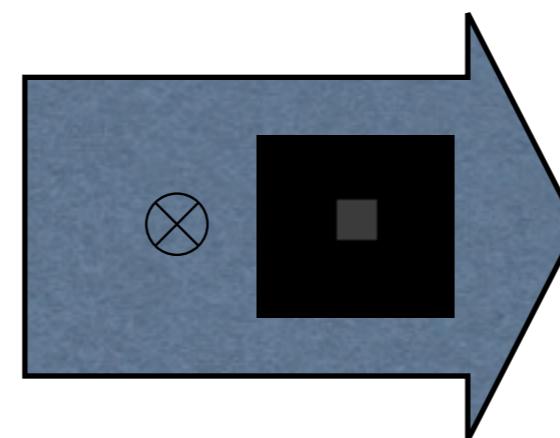
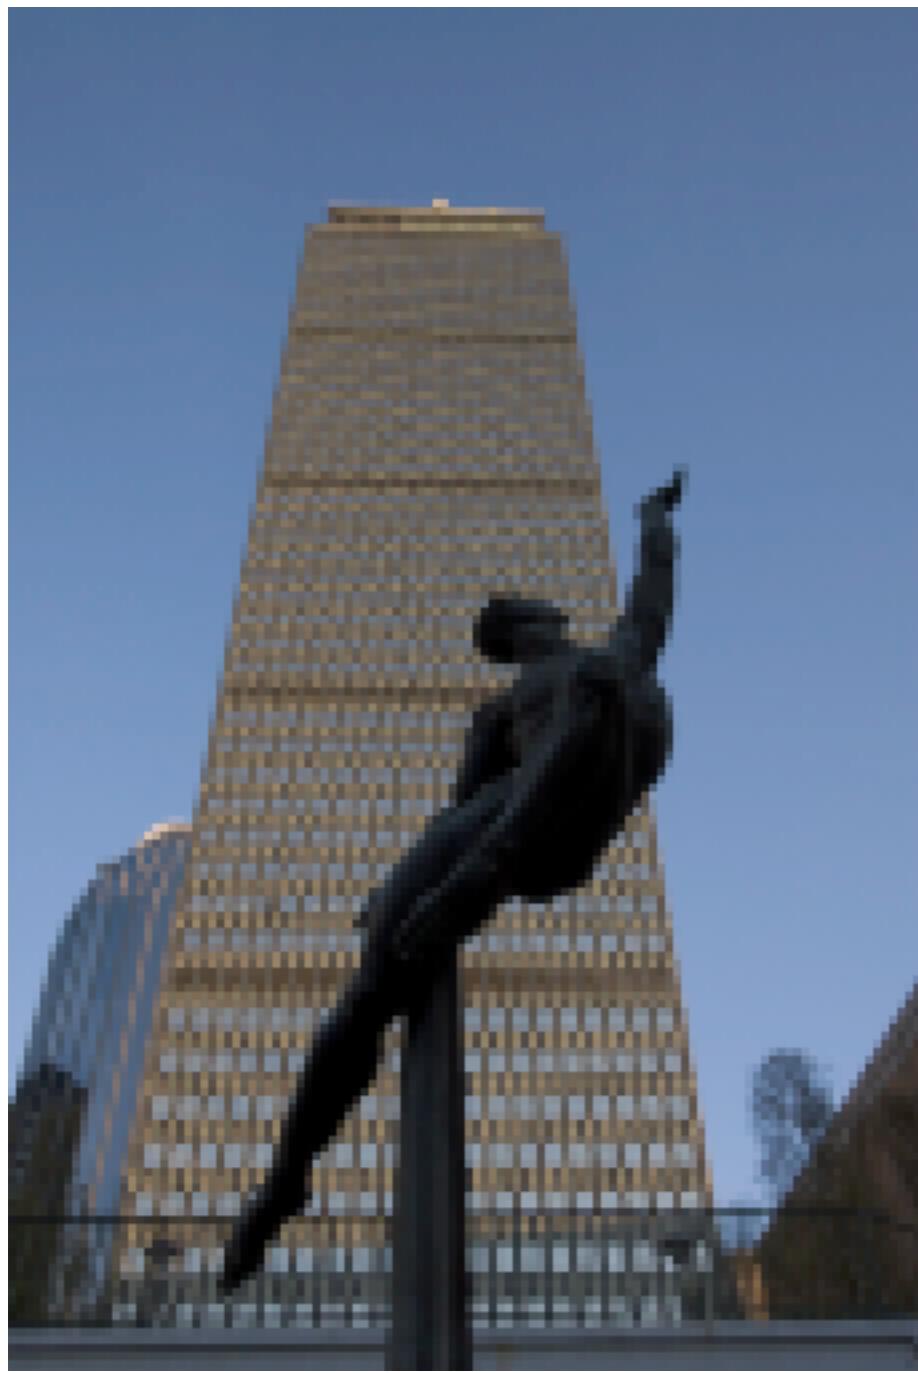
(b) Cone.

1	4	7	4	1
4	16	28	16	4
7	28	49	28	7
4	16	28	16	4
1	4	7	4	1

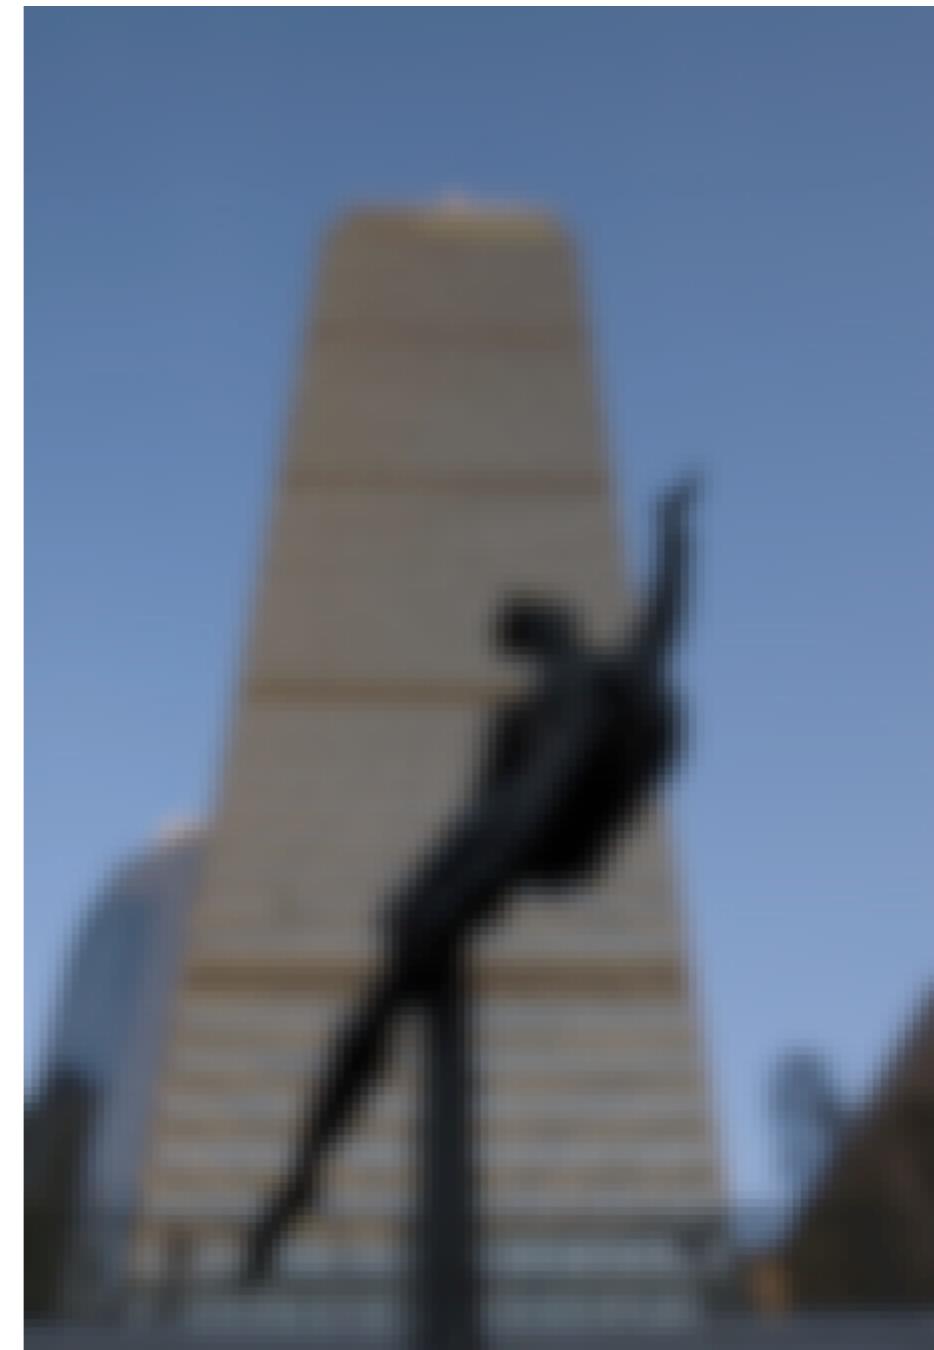
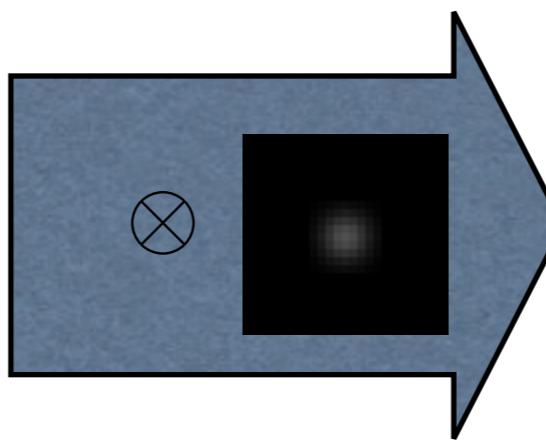
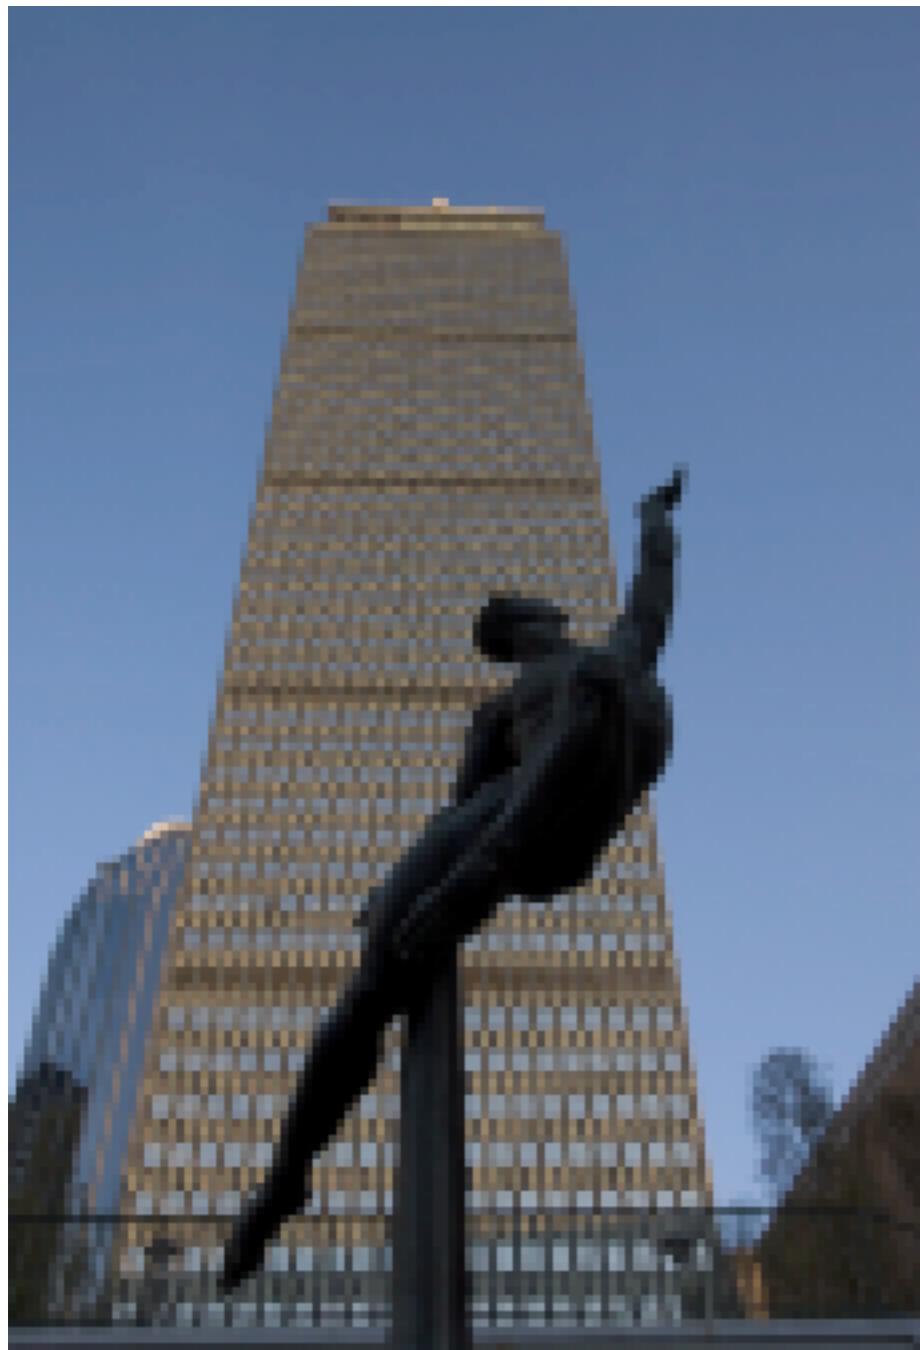
(c) Gaussian.

Table 6.1. Discretized kernels.

# Box Filter



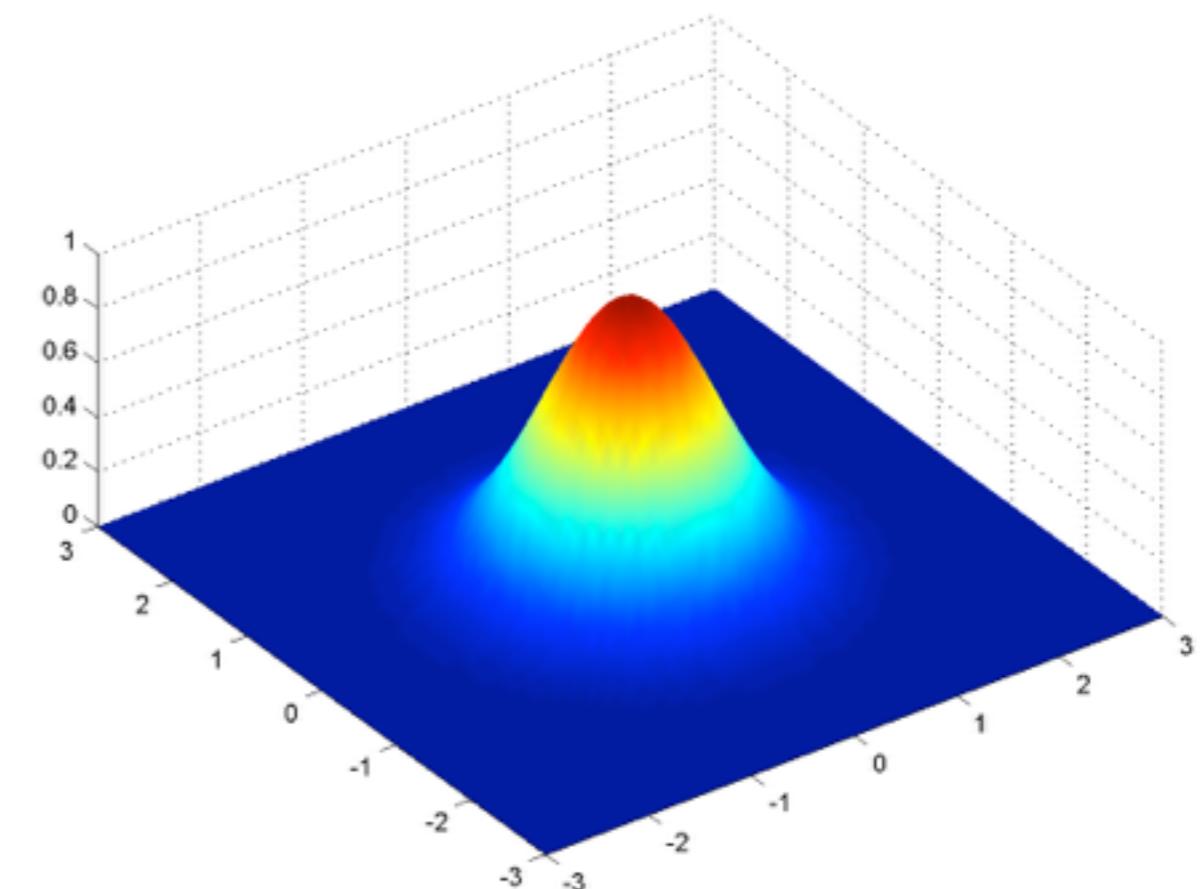
# Gaussian Filter



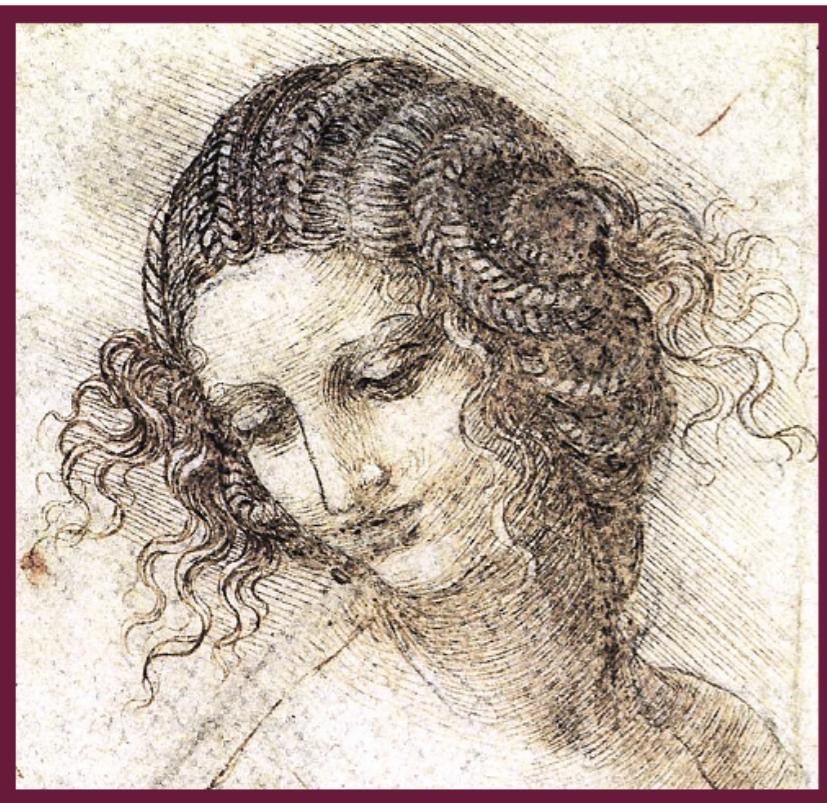
# Gaussians

$$G(x, y) = \frac{1}{2\pi\sigma^2} e^{\frac{-(x^2+y^2)}{2\sigma^2}}$$

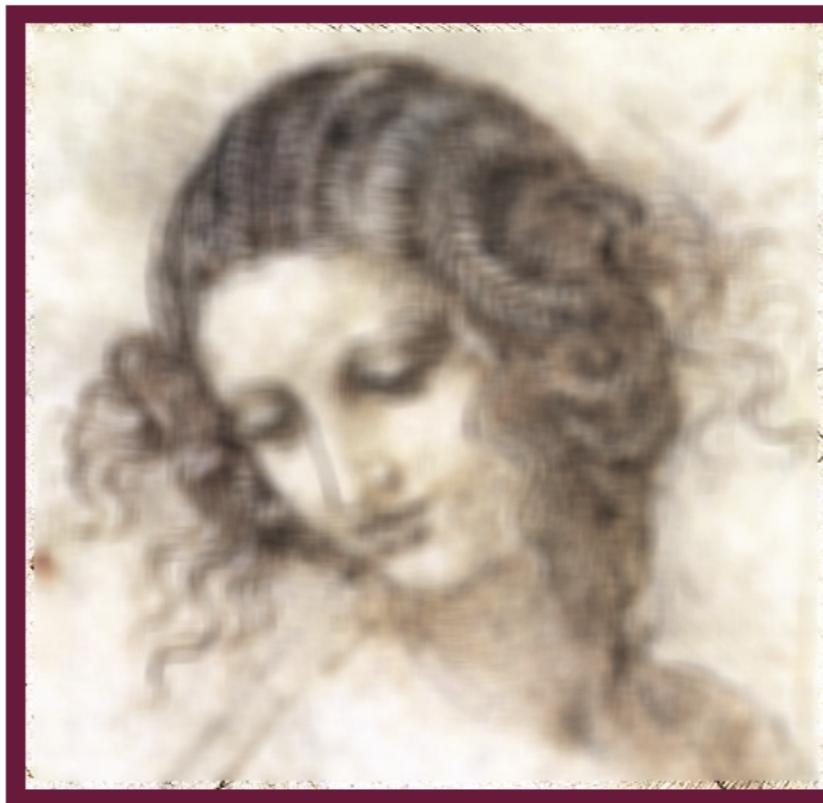
- Gaussian kernel is parameterized on the standard deviation  $\sigma$
- Large  $\sigma$ 's reduce the center peak and spread the information across a larger area
- Smaller  $\sigma$ 's create a thinner and taller peak
- Gaussians are smooth everywhere.
- Gaussians have infinite **support**
  - $>0$  everywhere
- But often truncate to  $2\sigma$  or  $3\sigma$



# Smoothing Example



(a) Source image.



(b)  $17 \times 17$  Box.

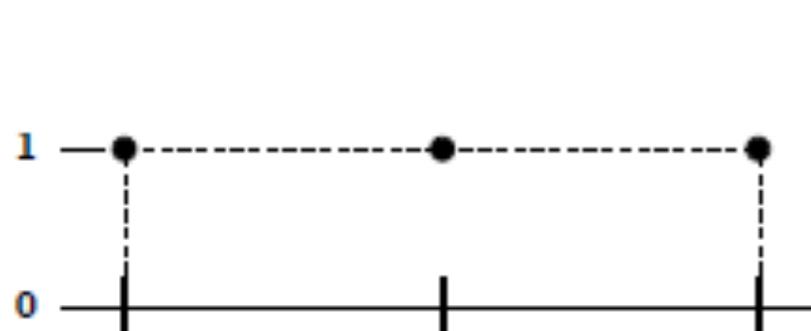


(c)  $17 \times 17$  Gaussian.

Figure 6.10. Smoothing examples.

# Smoothing Smoothing Filters

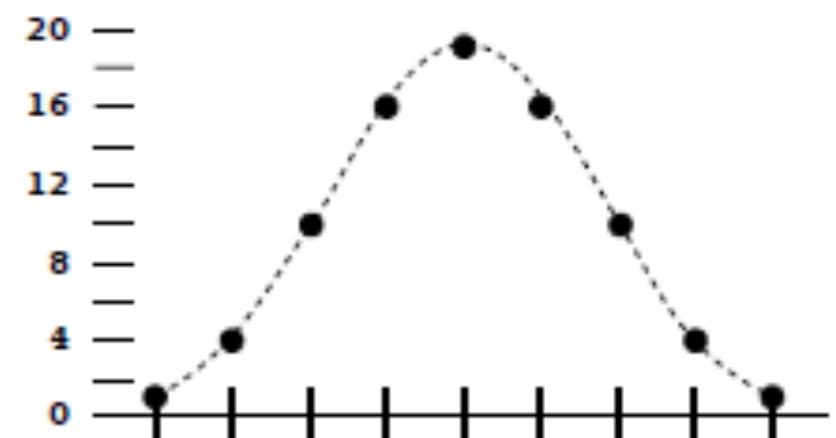
- Box  $\otimes$  Box = Tent (Pyramid)
- Tent  $\otimes$  Tent = Bell (Gaussian)



a) box filter



b) tent filter



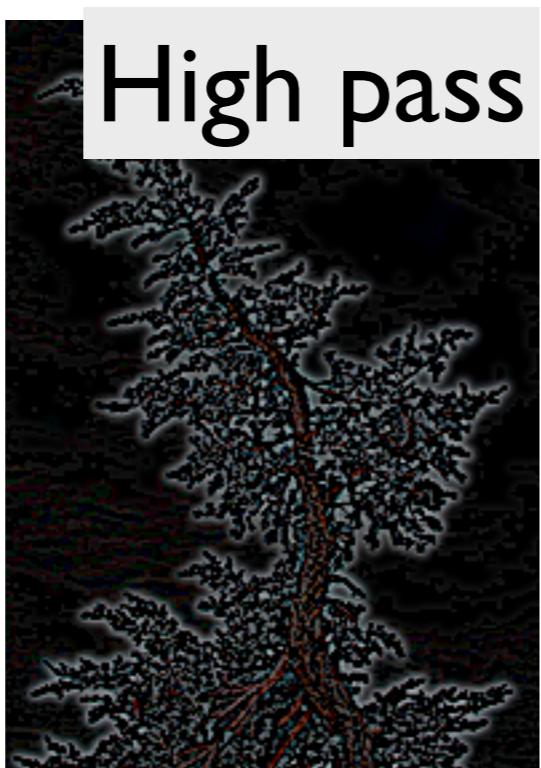
c) bell filter

# Sharpening Filters

# Sharpening (Idea)



High pass



Sharpened  
image

# Sharpening is a Convolution

- This equation can then be used to generate appropriate sharpening kernels
  - Assume that  $I = I \otimes K_{\text{identity}}$  and  $I_{\text{low}} = I \otimes K_{\text{low}}$

$$\begin{aligned}I_{\text{sharp}} &= (1 + \alpha)I - \alpha I_{\text{low}}, \\&= (1 + \alpha)(I \otimes K_{\text{identity}}) - \alpha(I \otimes K_{\text{low}}), \\&= I \otimes (1 + \alpha)K_{\text{identity}} - I \otimes (\alpha K_{\text{low}}), \\&= I \otimes ((1 + \alpha)K_{\text{identity}} - \alpha K_{\text{low}}).\end{aligned}$$

# Sharpening is a Convolution

$$\begin{aligned}I_{\text{sharp}} &= (1 + \alpha)I - \alpha I_{\text{low}}, \\&= (1 + \alpha)(I \otimes K_{\text{identity}}) - \alpha(I \otimes K_{\text{low}}), \\&= I \otimes (1 + \alpha)K_{\text{identity}} - I \otimes (\alpha K_{\text{low}}), \\&= I \otimes ((1 + \alpha)K_{\text{identity}} - \alpha K_{\text{low}}).\end{aligned}$$

$$K_{\text{identity}} = \frac{1}{9} \times \begin{bmatrix} 0 & 0 & 0 \\ 0 & 9 & 0 \\ 0 & 0 & 0 \end{bmatrix},$$

$$K_{\text{low}} = \frac{1}{9} \times \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix},$$

$$(1 + \alpha)K_{\text{identity}} - \alpha K_{\text{low}} = \frac{1}{9} \times \begin{bmatrix} -\alpha & -\alpha & -\alpha \\ -\alpha & (9 + 8\alpha) & -\alpha \\ -\alpha & -\alpha & -\alpha \end{bmatrix}$$

# Lec13 Required Reading

- Hunt — 3.4, 7.1.3, 7.1.4
- Szeliski — 2.3, 3.5:

<http://link.springer.com/book/10.1007/978-1-84882-935-0/page/1>